

DRAFT

January 9, 2013

Buffalo: An Aspect Oriented Programming Framework for C#

by

Wei Liao

A Project Report Submitted
in
Partial Fulfillment of the
Requirements for the Degree of
Master of Science
in
Computer Science

Supervised by

Prof. James E. Heliotis

Department of Computer Science

B. Thomas Golisano College of Computing and Information Sciences
Rochester Institute of Technology
Rochester, New York

February 2013

The project “Buffalo: An Aspect Oriented Programming Framework for C#” by Wei Liao has been examined and approved by the following Examination Committee:

Prof. James E. Heliotis
Professor
Project Committee Chair

Prof. Matthew Fluet
Associate Professor

Prof. Fereydoun Kazemian
Assistant Professor

Dedication

To Jackson and Evan

Acknowledgments

I am grateful for my adviser Prof. Heliotis, whose insightful advices; guidance and support from the beginning not only enabled me to complete the project on time, but also to a better understanding of the subject area.

I am also grateful for Prof. Fluet and Prof. Kazemian for their invaluable feed backs.

Last but not the least; I want to thank my wife Michelle, for all the support and encouragement during my years at school.

Abstract

Buffalo: An Aspect Oriented Programming Framework for C#

Wei Liao

Supervising Professor: Prof. James E. Heliotis

Aspect Oriented Programming (AOP) is a paradigm that let programmer isolate and separate crosscutting concerns from their programs. The concept has not been widely adopted by modern languages; support in tooling such as Integrated Development Environment (IDE) is also rare. In this project I designed and implemented Buffalo, an AOP framework to provide this capability for the .NET platform.

Buffalo performs Common Intermediate Language instruction set modification according to the aspects written by developer, with the help of the Mono Cecil library. Buffalo is .NET attribute based, which mean developers with existing .NET skills will have little or no learning curve to get started. Buffalo will help increase developer productivity in many areas such as unhandled exception catching, tracing and logging, etc.

Contents

Dedication	iii
Acknowledgments	iv
Abstract	v
1 Introduction	1
2 Design	2
2.1 Compiler Support	2
2.2 Run-time Interception	2
2.3 Post Compilation Weaving	3
2.4 A Buffalo Aspect	4
2.5 MethodBoundaryAspect	5
2.6 MethodAroundAspect	6
3 Implementation	9
3.1 How to Apply an Aspect	9
3.2 Aspect Interface	11
3.3 MethodArgs	13
3.4 Visual Studio Solution Structure	13
3.5 Implementation Overview	15
3.6 MethodBoundaryAspect Implementation Detail	15
3.7 MethodAroundAspect Implementation Detail	17
3.8 MethodArgs Implementation Detail	19
4 Analysis	20
5 Conclusions	22
5.1 Current Status	22

5.2	Future Work	23
5.3	Lessons Learned	23
	Bibliography	24
A	User Manual	26
A.1	Compiling	26
A.2	Simple Profiler	26
A.3	Transaction Database Example	29
A.4	Integrate With MS-Build System	31

List of Tables

4.1	Line counts	20
-----	-----------------------	----

List of Figures

2.1	AOP Framework Using Proxy Pattern	3
2.2	Buffalo Model	4
2.3	Overview of MethodAroundAspect	7
3.1	Logical Inclusion	10
3.2	Aspect Inheritance	11
3.3	Solution Structure	14
3.4	CIL Interception Points	17
A.1	Adding Buffalo.targets	32

Chapter 1

Introduction

Object Oriented Programming (OOP) languages have given programmers a lot of freedom in expressing themselves in Object Oriented Design. However, they are still lacking in some areas when it comes to particular software design decision such as cross-cutting concern [1].

In this project, a framework called "Buffalo" is designed and implemented to solve this type of problem on the .NET platform. Buffalo makes use of the .NET attribute system to weave aspect code to any targeted methods. The design and rationale of the framework is discussed in section 2. The implementation detail is given in section 3.

The result indicates that by using Buffalo, developers can separate cross-cutting concerns from the core of the program for easy maintenance, and ultimately be more productive. The analysis is discussed in section 4.

The report concludes in section 5 with the current project status. A set of planned future works is also discussed and what is learned from undergoing this project.

Buffalo comprises around 1,200 line of source codes. A user manual is included in Appendix A, which contains some usage examples on how the system works. Instructions is also included on how to integrate Buffalo with MS-Build.

Chapter 2

Design

2.1 Compiler Support

There are several broad approaches to implementing an AOP framework. The ideal approach would be to extend the compiler of the target language to provide built-in support, thus making AOP the first class citizen. However there are very few languages out there that take this approach, among the few are Delphi Prism [2] and AspectJ [3, 4].

Microsoft is currently in the "wait and see" mode regarding support of AOP development in the C# compiler [5]. Alternative compiler such as Mono C# [6] is open source, so technically anyone can build AOP support into it. While that would have been a fun challenge, that would have been a fairly big undertaking, and there is concern that the project might not be finished in the time frame wanted.

That leaves framework support as the other viable option. There are several implementation techniques to provide AOP capabilities [7, 8, 9] via framework.

2.2 Run-time Interception

Early on the implementation approaches were narrowed down to between Run-time Interception and Compile Time Weaving. As its name suggested, run-time interception operates while the program is in execution. It uses the proxy pattern where client communicate with the target object via a proxy, and aspects are injected to the proxy. This enables run-time behavior of the program to be modified. Figure 2.1 illustrates this process.

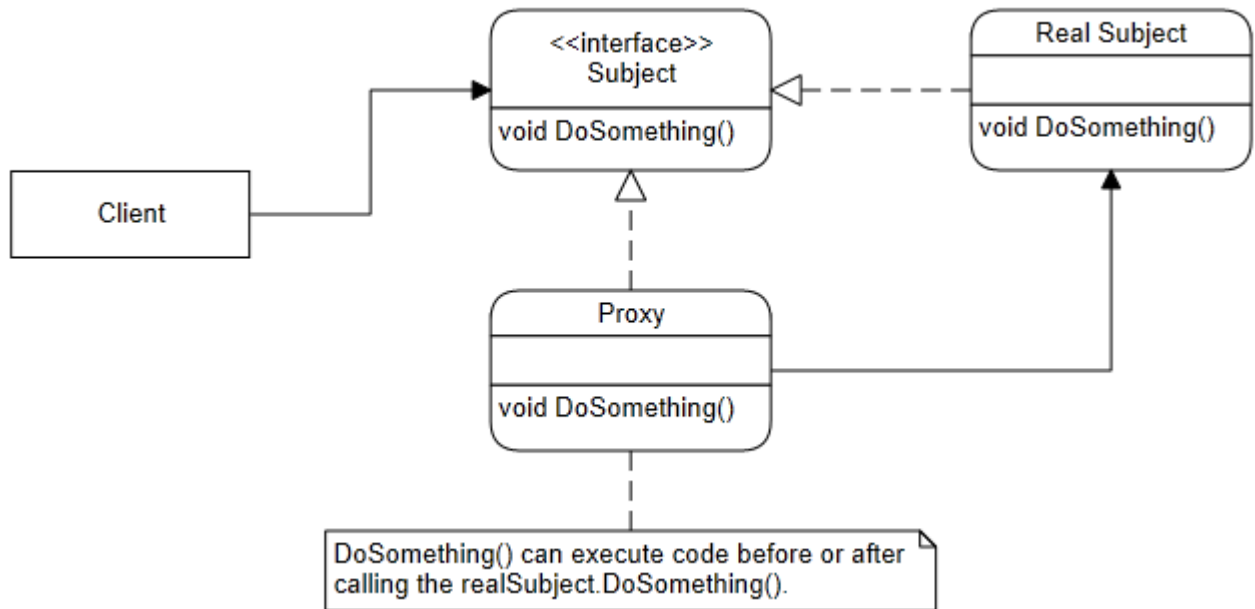


Figure 2.1: AOP Framework Using Proxy Pattern

New functionality can be added to the target object via the proxy. The disadvantage of this approach is that it involves the generation of proxy object at run-time. The run-time performance of the application will be impacted as the result. It is also restricting in that both target object and the proxy must implement a common interface for this to work, and that only virtual methods are exposed for interception.

From the end user's perspective, to use it the developer usually have to provide some type of mapping between the target object and the proxy via a configuration file so the actual proxy generation can occur. This approach although is easier to implement, but not as easy and friendly to use. Buffalo is not taking this approach mainly because one of the goal is to be flexible and simple to use.

2.3 Post Compilation Weaving

The approach Buffalo takes is Post Compilation Weaving. The idea is that after compilation of the source code, the framework takes over and disassembles the assembly. Buffalo

then weaves in the defined aspect code to all targeted methods. This approach is more difficult to implement as it involves modifying the underlying assembly by changing Common Intermediate Language (CIL) instructions [10]. But the advantage is that no run-time performance of proxy generation will be needed.

Since injection happens post-compilation, the whole process can be integrated into the MS-Build system to have the weaving invoked automatically if needed. This will further reduce the steps needed from the developer.

Figure 2.2 shows an overview of the compilation process and where Buffalo will fit in.

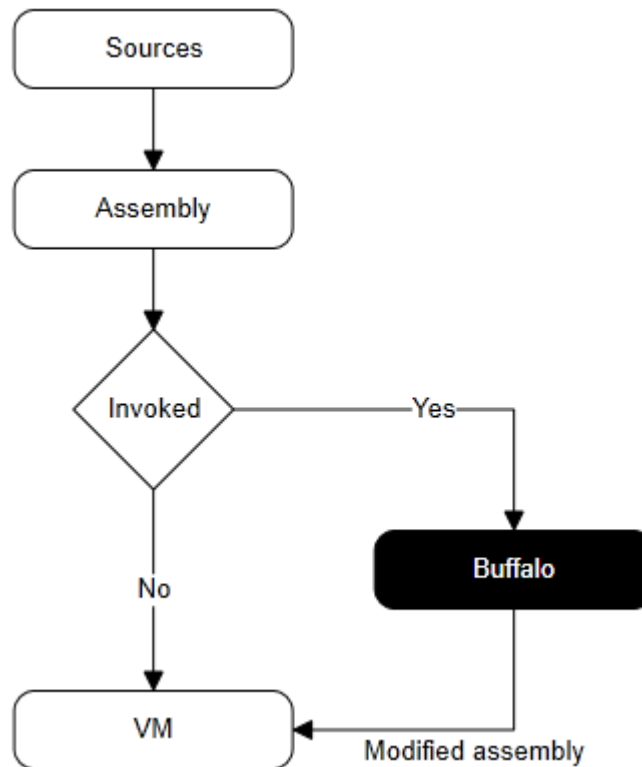


Figure 2.2: Buffalo Model

2.4 A Buffalo Aspect

When performing post compilation weaving, Buffalo has to be able to discover what aspect is applied to what methods in an assembly. In order to achieve that, the target assembly has

to carry some identifying meta-data.

A given .NET assembly already carries a great deal of such meta-data for various purpose. .NET has the `System.Attribute` type that exists primarily for the purpose of inserting meta-data into the assembly during compilation. When the source code is compiled, it is converted into CIL [11] and put inside a portable executable (PE) file, with the meta-data generated by the compiler.

Buffalo takes advantage of this characteristic in two phases.

1. An aspect defined in Buffalo will be in the form of an attribute, by sub-classing `System.Attribute`. It can contain any valid .NET code. But specifically an aspect needs to override various predefined methods in order to do something useful. The next section will discuss the relationship between various aspect types.
2. After compilation, the assembly will now contain the meta-data about the aspect. Buffalo can inspect the assembly for the information, and perform CIL code injection accordingly.

In other word, a Buffalo aspect is a .NET attribute in disguise.

2.5 MethodBoundaryAspect

What functionality does Buffalo support? What type of weaving does it do? For inspiration existing works such as AspectJ [3] and PostSharp [8] were studied. Specifically Buffalo will intercept the various point of an executing method. Those points are namely: before a method executes; after a method executes; whether or not the method executed successfully without error; or whether the method throws an exception any point during the execution. These various points of interception are grouped into the `MethodBoundaryAspect`.

`MethodBoundaryAspect` can be cleanly mapped to the try-catch-finally statements of the .NET languages. As far as the runtime is concern [12, 13], try-catch can be used liberally without serious performance degradation. For example, a simple method shown in Figure 2.1:

```
1 public void SomeFunction () {  
2     //Perform some action...  
3 }
```

Listing 2.1: Sample function

The above can be transformed by Buffalo into something shown in Figure 2.2. This clearly captures the spirit of the MethodBoundaryAspect. Regardless of whether the source already contain its own try-catch, or try-catch-finally blocks, the body of the method will be wrapped inside of a new try-catch-finally block by Buffalo.

```
1 public void SomeFunction () {  
2     try {  
3         OnBefore();  
4         //Perform some action  
5         OnSuccess();  
6     }  
7     catch (Exception e) {  
8         OnException(e);  
9     }  
10    finally {  
11        OnAfter();  
12    }  
13 }
```

Listing 2.2: Sample try-catch-finally

Transformed method in Figure 2.2 still does what the original method intends to do, only now at various point execution are being intercepted to provide more functionality.

2.6 MethodAroundAspect

Another type of aspect that Buffalo supports is the MethodAroundAspect. Rather than intercepting various execution points of a method, the method can be completely replaced by another method defined in an aspect, while preserving the option to call back into the original method if necessary.

At first glance `MethodAroundAspect` sounds straightforward to do, but it turns out to be much more involved than the `MethodBoundaryAspect`.

Since the option to call back into the original method is preserved, it is critical that under no circumstance should the original method be modified. If the method body instructions are simply overridden with that of the replacement, the call back to the original method will be meaningless since the method is now changed. The original method must be intact for the call back to happen.

To get around this obstacle, whenever Buffalo encounters the `MethodAroundAspect` applied to a method, it dynamically generates a replacement method in CIL with the same method signature as the original.

The body of this replacement method is also completely different than the original. It instantiates the aspect and makes call to the `Invoke()` method, which is the actual code that will be ran as a replacement.

Inside the `Invoke` method, developer can make a call back to the original method via a call to the `Proceed()` method. Then throughout the program, for any calls made to the original method, Buffalo would change them to call the replacement method instead. This is illustrated in figure 2.3.

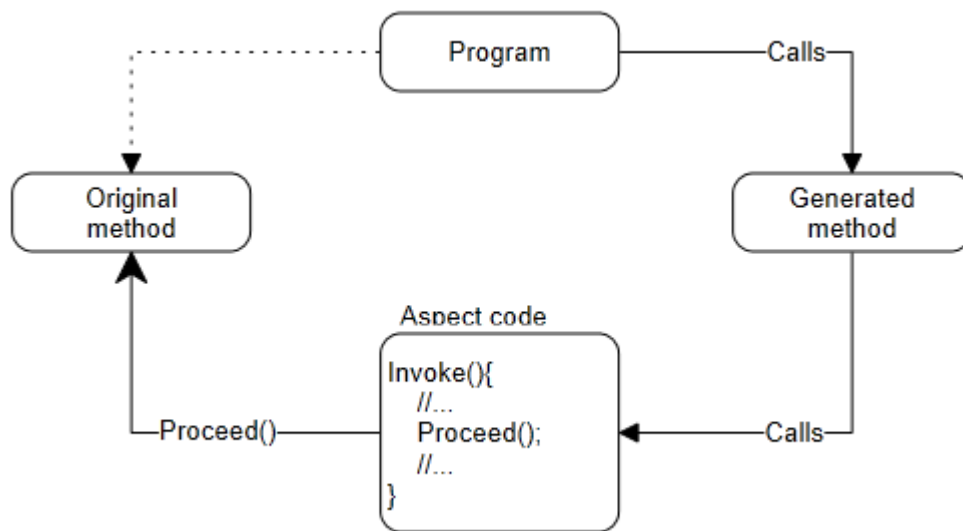


Figure 2.3: Overview of `MethodAroundAspect`

The dotted line from Program to the Original method indicates that once the MethodAroundAspect is applied to it, from the perspective of CIL the program cannot directly access that method any more. Access to the original method now has to come from inside the aspect. Also note that the original method is not changed at any given time.

Chapter 3

Implementation

3.1 How to Apply an Aspect

Since an aspect is really a .NET attribute, it can be used just like any other attribute. But code annotated with an aspect is special in that it can be understood only by Buffalo.

A Buffalo aspect can be applied on three levels, with the following characteristics:

1. Method - apply the aspect to an individual method.
2. Class - if aspect is applied to a class, all public methods including the public properties automatically get applied.
3. Assembly - if aspect is applied to an assembly, #2 will apply but for all the public classes within the assembly.

An exception to the above rule is the `MethodAroundAspect`, where it can only be applied on a method level, as will be shown later on.

All aspects have a property named `AttributeExclude`, if set to true then the annotated target will not be included in the weaving. This exclusion can happen on any levels. For example, if a method contains this annotation `[SampleAspect(AttributeExclude=true)]`, the method will be skipped for the `SampleAspect` during the weaving process.

No matter how the aspect is applied, ultimately it will result in a list of the methods that are annotated. This simply mean if the aspect is applied to a single method, that method is the only one that will get CIL modified. If the aspect is applied on the whole assembly, then all public methods will be CIL modified.

To get the list of the eligible methods for CIL modification, Buffalo attempts various checking according to figure 3.1 to see if it should include a given method.

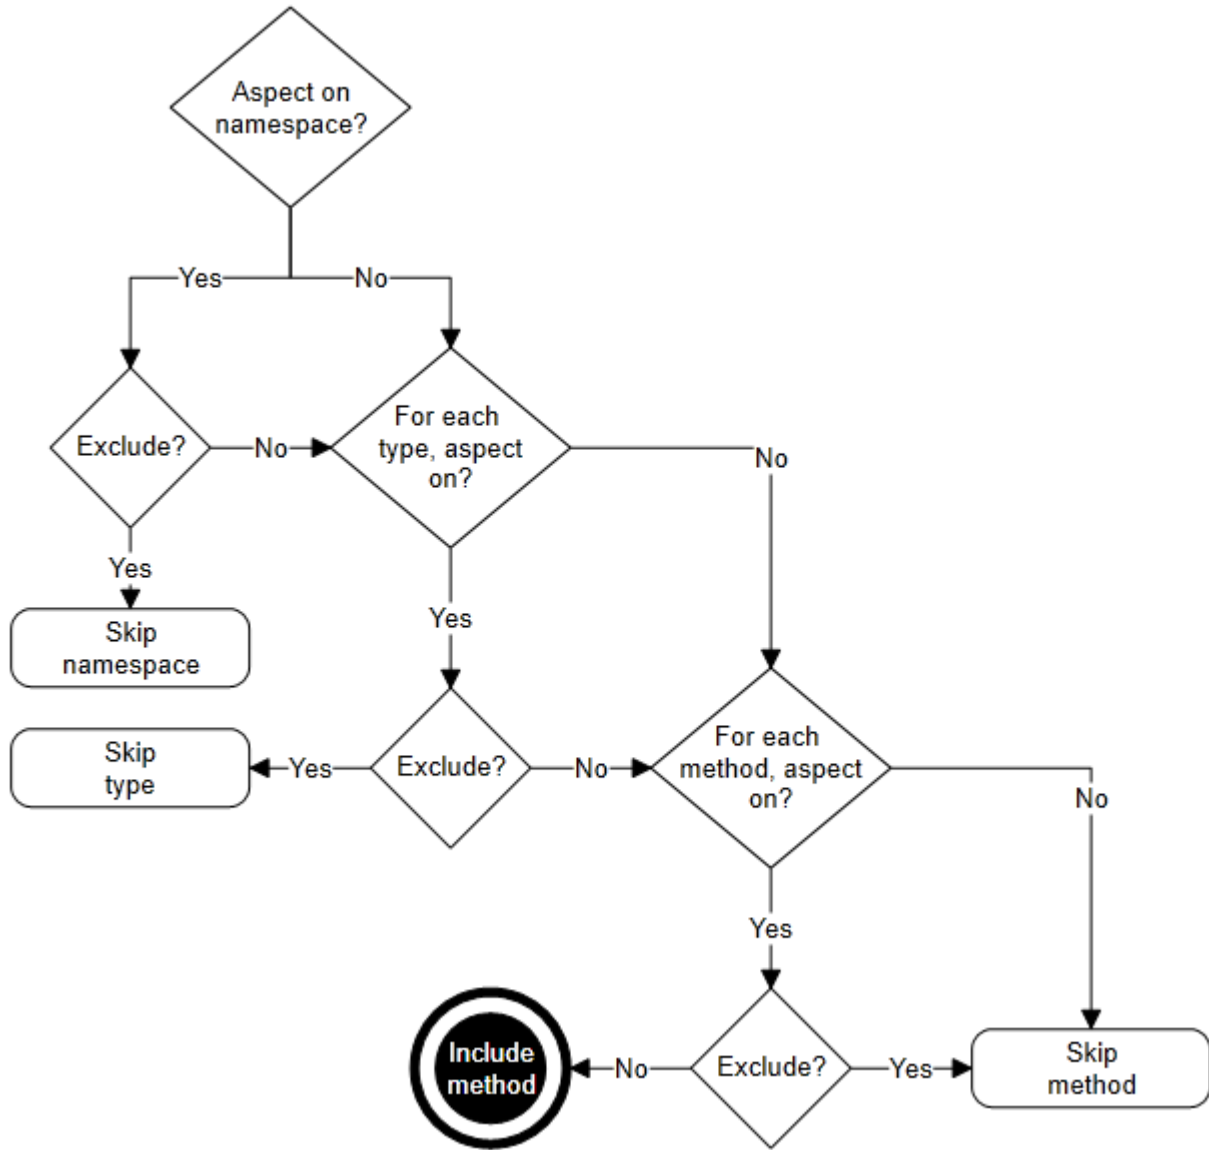


Figure 3.1: Logical Inclusion

If no aspect is applied on the assembly, that does not necessarily mean no aspect is applied anywhere, the aspect might still be applied on any given class or method.

The take away from the above diagram, is that Buffalo first checks if an aspect is applied to the target, then check if it is set to be excluded. At the end it will end up with a list of

methods that should be CIL modified.

3.2 Aspect Interface

Figure 3.2 shows the relationship of various aspect types in Buffalo. This is used by Buffalo to identify aspects during reflection.

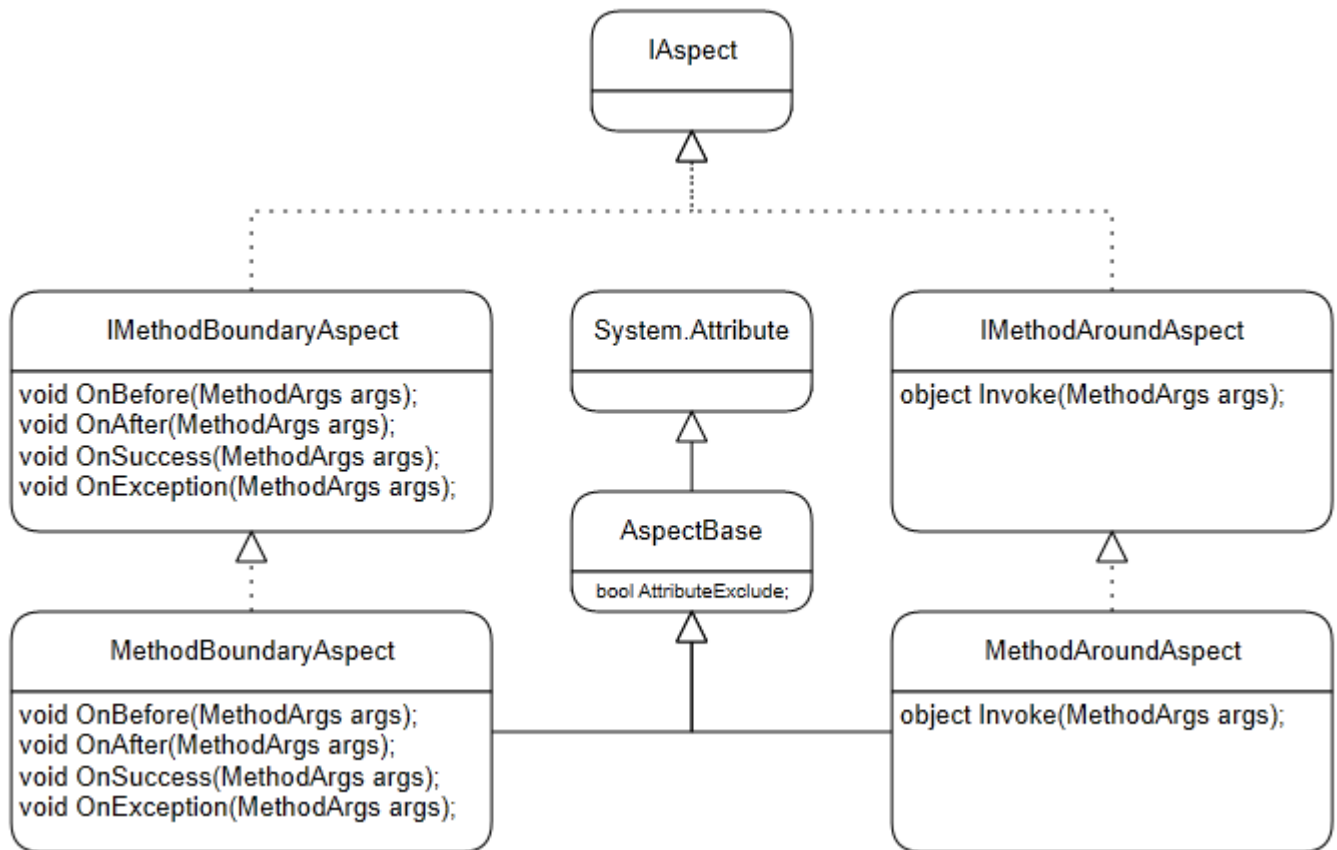


Figure 3.2: Aspect Inheritance

All aspects ultimately implements the IAspect interface, therefore it can be reasoned that for all the public types in an assembly, if it implements IAspect, then it must be an aspect itself.

Buffalo supports more than one aspect applied at any given level. This will allow developers more flexibility while developing multiple aspects and applying them as needed.

Furthermore, by default, an aspect will be automatically excluded from applying to itself. This is implemented to prevent stack overflow in some cases. Although argument can be made that an aspect should be able to be applied to a different aspect; that is not currently implemented in Buffalo.

Listing 3.1 shows how an aspect is created. It inherits `MethodBoundaryAspect` and overrides `OnBefore` and `OnAfter`.

```
1 using Buffalo;
2 using System;
3
4 public class TraceAspect : MethodBoundaryAspect
5 {
6     public override void OnBefore(MethodArgs args)
7     {
8         Display("ENTERING", args);
9     }
10
11    public override void OnAfter(MethodArgs args)
12    {
13        Display("EXITING", args);
14    }
15
16    void Display(string title, MethodArgs args)
17    {
18        Console.WriteLine("{0} {1}", title, args.FullName);
19        foreach (var p in args.Parameters)
20        {
21            Console.WriteLine("\t{0} ({1}) = {2}",
22                               p.Name, p.Type, p.Value);
23        }
24    }
25 }
```

Listing 3.1: Sample TraceAspect

To use the aspect, simply apply it to a method, class or assembly. Once the code is compiled to produce an assembly, `BuffaloAOP.exe` can be invoked by passing it the path to the assembly. Buffalo will take over and weave in the aspect. This and more examples and details are provided in appendix A.

```
1 [TraceAspect]
2 public class Hello
3 {
4     //...
5 }
```

Listing 3.2: Apply Aspect on Class Level

3.3 MethodArgs

As mentioned above, when all is said and done, an aspect ultimately gets injected into each *individual* method. When developing an aspect, meta-information about the target method can be accessed. This information is encapsulated via the MethodArgs object passed in as parameter to the aspect. Currently the full method signature, the method name, return type and parameter list including parameter name, type and value are captured for each target method.

The parameter list capturing is especially of interest, it enables developer to peek inside the method that is executing at various point and inspect its parameter values. This will be useful in case such as exception handling, where it will be useful to actually see what the values were at the time of the exception.

3.4 Visual Studio Solution Structure

Originally Buffalo was implemented as one executable; that includes the various aspects and the program that initiates the weaving. It was later on separated into two assemblies. One is a class library that contains the actual implementation. Another is a command line executable that calls into the class library to perform the weaving. This separation is necessary so developer can perform weaving from the command line or hook into MS-Build if necessary.

To actually write the aspect, one only needs to reference the class library as underlined in figure 3.3, which is much cleaner than referencing an executable.

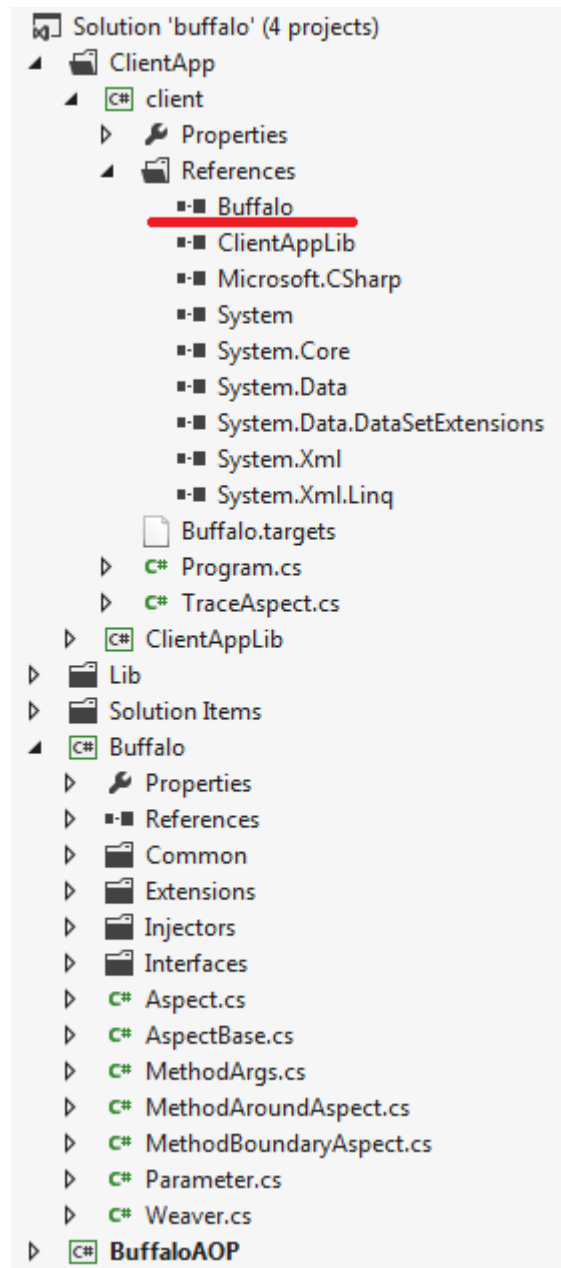


Figure 3.3: Solution Structure

The client project shown above is a simple program included in the solution for testing.

3.5 Implementation Overview

The implementation begins by finding all eligible methods to be injected using the verifying process indicated in figure 3.1. Each eligible method will have one or more aspects applied to it. The actual injection process will loop through each aspects for the eligible method and inject the necessary CIL instructions.

The CIL instructions modification is performed by the open source Mono.Cecil library. Originally the System.Reflection APIs in the .NET Framework was favored, but it was later determined that Mono.Cecil provides more features and is much easier to work with. Buffalo uses Mono.Cecil heavily to modify CIL instructions and to assemble the final assembly.

3.6 MethodBoundaryAspect Implementation Detail

Each type of aspect has its own injector that implements the IInjectable interface. This interface contains only one method contract - Inject(..). It takes the list of eligible methods and injects the appropriate aspect to them.

MethodBoundaryAspect is pretty straightforward to implement. Take the following hello world example, wrapped in a try-catch-finally block as mentioned previously:

```
1 public void SayHello()  
2 {  
3     try{  
4         Console.WriteLine(Hello World!);  
5     }catch(Exception ex){  
6     }finally{  
7     }  
8 }
```

Listing 3.3: SayHello function

The generated CIL is shown in figure 3.4. For ease of display the CIL has been cleaned up a bit:

```
1 .try
```



```
2 {  
3     .try  
4     {  
5         IL_0002: Ldstr "Hello World!"  
6         IL_0007: call void [mscorlib]System.Console::WriteLine(string)  
7         IL_000e: leave.s IL0015  
8     }  
9     catch [mscorlib]System.Exception  
10    {  
11        IL_0010: stloc.0  
12        IL_0013: leave.s IL_0015  
13    }  
14    IL_0015: leave.s IL_001c  
15 }  
16 finally  
17 {  
18     IL_001a: endfinally  
19 }  
20 IL_001c: ret
```

Listing 3.4: CIL generated for sample C# function

Figure 3.4 shows the standard emission of the CLR when it encounters the try-catch-finally statement. In CLR there is a concept of the protected region, where each region is associated with a handler. A try-catch-finally is actually encapsulated in two such regions: a catch and a finally. From here it can be easily figured out where to inject the various boundary aspects, as shown in figure 3.4.




Figure 3.4: CIL Interception Points

3.7 MethodAroundAspect Implementation Detail

The Around aspect on the other hand is a much more complicated compared to the Boundary aspect.

MethodAroundAspect implements `IMethodAroundAspect` which has the following contract:

```

1 internal interface IMethodAroundAspect : IAspect
2 {
3     void Invoke(MethodArgs args)
4 }
  
```

Listing 3.5: IMethodAroundAspect Contract

When developing an aspect a `Proceed()` can be issued to signal a call back into the original method. The steps taken to implement `MethodAroundAspect` in CIL are roughly

as follow:

1. Create a replacement for the annotated function with exactly the same method signature.
2. Create and store a variable pointing to the aspect.
3. Copy all parameters from original method to the newly created replacement function.
4. Create a variable to hold MethodArgs.
5. Issue a call to Invoke() from the replacement function, passing in the MethodArgs variable.
6. Handle the return value appropriately.
 - (a) If original method returns non void type, put the return value back on the stack.
 - (b) If original method returns void, we need to discard the return value from Invoke().
7. Handle Proceed() that might be issued from inside the Invoke().
 - (a) Load all the parameters onto the stack.
 - (b) Call back into the original method.
 - (c) Handle the return value appropriately.
8. Modify all calls from original method to the replacement method.

As figure 2.3 shown, the actual calling of either the original or replacement method is abstracted away. This is also a testament of the saying in Software Engineering that "anything can be resolved by another layer of abstraction".

Another important distinction is that MethodAroundAspect currently can be applied only on the method level, and that it should be applied to one method only. This is by design because a replacement method might not be appropriate to replace more than one

method. Especially if it is applied on the assembly level, all the public methods will be replaced by a single replacement method!

3.8 MethodArgs Implementation Detail

When developing an aspect, information about the target method can be accessed. To achieve that `MethodArgs` is used. This is the object passed into each aspect. During the weaving, an instance of `MethodArgs` is created, with all properties assembled dynamically to capture the information of the current executing method. `MethodArgs` is then passed as parameter into each of the aspect.

Being able to capture some information about the annotated methods will be useful. For example, in case of a profiling aspect, information about the method at the time it was access will be helpful. Being able to look at the parameter values in case of error will also be extremely useful in case of debugging.

A distinct instance of `MethodArgs` for each boundary aspects was instantiated at an early Buffalo implementation. Later on as an optimization only one instance is instantiated at the beginning of the method body and that instance is used in all the boundary aspects for a target method.

An example of how to use `MethodArgs` is presented in the user manual.

Chapter 4

Analysis

The project hypothesized that by using Buffalo, programmers can separate the cross-cutting concerns from their applications quickly and easily. Since the concerns are encapsulated in a distinct unit of code, it also enables programmers to easily maintain the aspects and modify them as needed.

One analysis performed is to write an aspect to catch unhandled exceptions in test programs. The size of the test programs varies from comprising of 50 methods to 1,000 methods. Suppose that to manually implement the exception handling, a programmer will have to write on average 5 lines of code to catch the exception.

Programs	Lines (Traditions)	Lines (Buffalo)
50	250	0-1
500	2,500	0-1
1,000	5,000	0-1

Table 4.1: Line counts

If exception handling is implemented for every method by hand, more line of the same try-catch block of code will have to be written as the application adds more methods. The number of line of repetitive code would increase linearly.

Lines of code have a direct correlation to the cost of the development as it will take programmers more time. And this will also have a direct impact on application release schedule.

By using a framework like Buffalo, unhandled exception can be centralized in one aspect, and then simply apply it to every method by applying it on the assembly level.

As a result the source code is free from the repetitive try-catch-finally blocks. The line of code we have to write is one line at most, and will stay constant even as more types and methods are added to the application. This will also give developer a peace of mind that every method will be handled automatically.

One can argue that since unhandled exceptions will bubble up the chain, a developer can simply catch them in the main method, and this would have achieved similar effect. That approach is limited in that it is very generic. When the main method catches an exception, it has no idea what the internal state of the method that failed was. Using Buffalo, the internal state can be inspected by checking the MethodArgs object; this will help in case of debugging.

Buffalo allows developer to quickly create aspects to solve various problems, from unhandled exception to instrumenting the application. These are just some of the scenarios where Buffalo proves to be helpful. More examples are provided in appendix A.

Buffalo has performed well in isolating cross-cutting concerns into single unit of code, which is easily maintained and modified.

Chapter 5

Conclusions

5.1 Current Status

Buffalo is currently at version 0.2. It contains two types of aspect: `MethodBoundaryAspect` - where various execution points can be intercepted. `MethodAroundAspect` - where a method can be completely replaced while preserving the option to call back into the original method. `MethodAroundAspect` targets individual method, whereas `MethodBoundaryAspect` can be applied on three levels on a .NET application.

Originally MSBuild integration was planned, that when the Buffalo program was installed via the setup executable, it will modify the relevant .NET configuration files to hook `BuffaloAOP.exe` into MSBuild. That would trigger the weaving process from within the Visual Studio IDE when the solution is compiled. It was later found that in the latest version of Visual Studio 2012, Microsoft has dropped support for the setup project type. As a work around, instruction is provided in appendix A to manually hook into MSBuild.

Originally `MethodAroundAspect` was intended to be used similar to `MethodBoundaryAspect`, where it can be applied on three levels. However it was later determined that since the CIL instruction of the actual aspect will also have to be modified, it really does not make sense for it to be applied to more than one method, as it would introduce conflicting changes.

A possible solution to `MethodAroundAspect` would be to clone the aspect if it is applied to more than one method. However the cloning process is not as easy to do since everything would have to be copied, including references to whatever objects used in the aspect.

5.2 Future Work

There are couple areas where Buffalo can be improved upon. Usability wise, currently there is no automatic setup program that installs Buffalo onto users computer. There is also no automatic integration into MS-Build System. Some manual steps are still needed in order to provide a more seamless experience.

Functionality wise, when Buffalo performs post compilation weaving it starts fresh each time; it would be interesting to see how incremental weaving can be done here. Another useful functionality is to be able to apply aspects by matching a set of methods.

MethodAroundAspect can be improved to be applied on three levels by possibly cloning the aspect.

5.3 Lessons Learned

A framework such as Buffalo mitigates the problem of cross-cutting concerns. Still, to efficiently tackle the root of the problem, compiler vendors have to actively embrace the AOP concept and provide native support in their languages.

Developers also have to understand such problems and what solutions are available to better educate themselves.

Only when the concept is widely understood and supported by both developers and vendors can there hope to begin alleviating such problems.

Bibliography

- [1] et al. Gregor Kiczales, John Lamping. Aspect-Oriented Programming. In *Proc. of European Conference on Object-Oriented Programming (ECOOP)*, Finland, 1997.
- [2] Cirrus - The Oxygene Language Wiki. <http://wiki.oxygenelanguage.com/en/Cirrus>. Accessed: 07/16/2012.
- [3] The Eclipse Foundation. AspectJ Frequently Asked Questions. <http://www.eclipse.org/aspectj/doc/released/faq.php>. Accessed: 07/19/2012.
- [4] Ramnivas Ladda. *Aspectj in Action: Enterprise AOP with Spring Applications*. Manning Publications, second edition, October 2009.
- [5] MSDN TV. Transcript: Whiteboard with Anders Hejlsberg. <http://msdn.microsoft.com/en-us/cc320411.aspx>. Accessed: 07/19/2012.
- [6] Mono. CSharp Compiler - Mono. http://www.mono-project.com/CSharp_Compiler. Accessed: 07/19/2012.
- [7] M. Devi Prasad and B. D. Chaudhary. AOP Support for C#. In *Proc. of Second International Conference on Aspect-Oriented Software Development*, pages 49–53, 2003.
- [8] Gael Fraitteur. User-friendly aspects with compile-time imperative semantics in .NET: an overview of PostSharp. In *International Conference on Aspect-Oriented Software Development*, 2008.
- [9] Howard Kim. AspectC#: An AOSD implementation for C#. Master’s thesis, Trinity College Dublin, 2002.
- [10] Aleksandr Mikunov. Rewrite MSIL Code on the Fly with the .NET Framework Proling API. *MSDN Magazine*, 2003.
- [11] Serge Lidin. *Expert .NET 2.0 IL Assembler*. Apress, first edition, August 2006.

- [12] ECMA International. ECMA-334 - C# Language Specification. <http://www.ecma-international.org/publications/files/ECMA-ST/Ecma-334.pdf>, June 2006.
- [13] ECMA International. ECMA-335 - Common Language Infrastructure (CLI). <http://www.ecma-international.org/publications/files/ECMA-ST/Ecma-335.pdf>, June 2012.

Appendix A

User Manual

It is easy to get start using Buffalo. I assume you are reasonably familiar with the Visual Studio IDE and the general layout of a VS solution. I also assume you know how to compile a solution and where to find the compiled assembly. Buffalo is developed in VS2012, it is recommended you have the same version of the IDE installed.

A.1 Compiling

To begin, first download the full source code from <https://github.com/wliao008/buffalo>. Open buffalo.sln in VS2012 and compile the source code. This will produce the Buffalo.dll and BuffaloAOP.exe in their respective bin/debug folder.

For this example we will perform the weaving from a command prompt. So create a folder name under the C drive, here I will call it "Buffalo". Copy Buffalo.dll, BuffaloAOP.exe and Mono.Cecil.dll to C:\Buffalo. Note this folder can be located anywhere in your system, I am just putting it on the C drive for simplicity.

Now let us create an aspect.

A.2 Simple Profiler

In this example we will create a profiler for our application. Suppose we have the following simple program.

```
1 using System;  
2  
3 namespace Hello
```

```
4 {
5     class Program
6     {
7         static void Main(string[] args)
8         {
9             Hello h = new Hello();
10            h.SayHello();
11            h.Say("Hey Buffalo how's it going!");
12
13            //pause the console
14            Console.Read();
15        }
16    }
17
18    public class Hello
19    {
20        public void SayHello()
21        {
22            Console.WriteLine("Hello World!");
23        }
24
25        public void Say(string msg)
26        {
27            Console.WriteLine(msg);
28        }
29    }
30 }
```

Listing A.1: Hello program

When the program runs, it will display the following output:

```
1 Hello World!
2 Hey Buffalo how's it going!
```

Listing A.2: Hello program output

And suppose that we want to monitor the program, we want to know when a method was accessed and exited. We can easily create a aspect to do such work.

```
1 using Buffalo;
2 using System;
3
4 public class TraceAspect : MethodBoundaryAspect
5 {
6     public override void OnBefore(MethodArgs args)
7     {
8         Display("ENTERING", args);
9     }
10
11    public override void OnAfter(MethodArgs args)
12    {
13        Display("EXITING", args);
14    }
15
16    public override void OnSuccess(MethodArgs args)
```

```

17     {
18         Display("SUCCESSFULLY EXECUTED", args);
19     }
20
21     public override void OnException(MethodArgs args)
22     {
23         Display("EXCEPTION ON", args);
24     }
25
26     void Display(string title, MethodArgs args)
27     {
28         Console.WriteLine("{0} {1}", title, args.FullName);
29         foreach (var p in args.Parameters)
30         {
31             Console.WriteLine("\t{0} ({1}) = {2}", p.Name, p.Type, p.Value);
32         }
33     }
34 }

```

Listing A.3: TraceAspect

With the aspect defined, now we can apply this aspect on any of the three different levels. Lets apply it to the Hello class for example.

```

1 [TraceAspect]
2 public class Hello
3 {
4     //...
5 }

```

Listing A.4: Apply Aspect to the Hello Class

Now everything is in place. We can now invoke the BuffaloAOP.exe to perform the weaving. Open a command prompt and navigate to C:\Buffalo. And issue this command:

```

1 C:\Buffalo>BuffaloAOP.exe <path_to_the_hello_program.exe>

```

Listing A.5: Invoking BuffaloAOP.exe

Replace path to the hello program exe with the actual complete path to the program assembly. Suppose the program assembly is located at C:\Projects\Hello\bin\Hello.exe, we would issue the command as follow:

```

1 C:\Buffalo>BuffaloAOP.exe C:\Projects\Hello\bin\Hello.exe

```

Listing A.6: Invoking BuffaloAOP.exe Example

If everything goes well BuffaloAOP.exe will perform the injection and put the final assembly in the Modified folder inside the folder of the target assembly. In this case it will

be at C:\Projects\Hello\bin\Modified\Hello.exe. Now when the program runs, it will display the following output:

```

1 ENTERING System.Void Hello.Program::Main(System.String[])
2     args (System.String[]) = System.String[]
3 ENTERING System.Void Hello.Hello::.ctor()
4 SUCCESSFULLY EXECUTED System.Void Hello.Hello::.ctor()
5 EXITING System.Void Hello.Hello::.ctor()
6 ENTERING System.Void Hello.Hello::SayHello()
7 Hello World!
8 SUCCESSFULLY EXECUTED System.Void Hello.Hello::SayHello()
9 EXITING System.Void Hello.Hello::SayHello()
10 ENTERING System.Void Hello.Hello::Say(System.String)
11     msg (System.String) = Hey Buffalo how's it going!
12 Hey Buffalo how's it going!
13 SUCCESSFULLY EXECUTED System.Void Hello.Hello::Say(System.String)
14     msg (System.String) = Hey Buffalo how's it going!
15 EXITING System.Void Hello.Hello::Say(System.String)
16     msg (System.String) = Hey Buffalo how's it going!

```

Listing A.7: TraceAspect output

Line 7 and 12 are the original method output, the rest are the output of the various interception points. Note that line 2, 11, 14 and 16 show the parameter value passed into each method.

A.3 Transaction Database Example

It is common to work with code that saves information to a database. And it is typical that the data might have to go into different tables. For example, let us look at an online ordering system. When a customer clicks the buy button during checkout, an order has to be created, and each order item has to be created as well. Order and order items usually go into different tables since they have the logical parent-child relationship. So we might have the following code to save the data:

```

1 public void SaveOrder(string customerName, string item1, string item2)
2 {
3     //create the linq-to-entity context
4     ModelContainer db = new ModelContainer();
5     //create an order obj
6     Order o = new Order();
7     o.CustomerName = customerName;
8     db.Orders.Add(o);
9     db.SaveChanges(); // <-save the order!
10
11     List<string> items = new List<string>();

```

```
12         if (!string.IsNullOrEmpty(item1)) items.Add(item1);
13         if (!string.IsNullOrEmpty(item2)) items.Add(item2);
14         items.ForEach(x =>
15         {
16             //create the order item
17             OrderItem item = new OrderItem();
18             item.ItemName = x;
19             item.OrderId = -1; //<-- intentionally failing the save since there is no such order id
20             db.OrderItems.Add(item);
21         });
22
23         db.SaveChanges(); //<-- saving the order items
24     }
```

Listing A.8: SaveOrder Example

The above first saves an order, and then saves each individual order items. The data goes into different tables. Note that we are using the Linq to Entity ORM here, technically it handles transaction if `db.SaveChanges()` at line 9 is removed and is calls at the end with one call. This code is for illustration purpose only. If you are not using Linq to Entity ORM, or sometime is not uncommon to issue multiple `db.SaveChanges()` throughout the source.

Also note that at line 19 we are intentionally failing the `OrderItem` save. At that point the `Order` itself has been saved. When `OrderItem` fails, we ended up with incomplete data in the database.

This case such as this, it is important that the atomicity rule is enforced. The save operation should either be completely saved or nothing should be saved at all.

Atomicity rule is usually enforced by using transaction. So we can modify the above code like this:

```
1 public void SaveOrderManualTransaction(string customerName, string item1, string item2)
2 {
3     using (TransactionScope scope = new TransactionScope())
4     {
5         //same body of code to save..
6     }
7 }
```

Listing A.9: SaveOrderManualTransaction Example

The save operations are now wrapped in a `TransactionScope` object. If anything bad happens on any of the `db.SaveChanges()` calls the whole operation is aborted, so we wont

ended up with incomplete data in the database.

Scenario like this is a suitable candidate for using aspect. If we need to enforce atomicity elsewhere, we would have to do something repetitive with `TransactionScope`. We can refactor and extract out the aspect using Buffalo.

```
1 public class RunInTransactionAspect : MethodBoundaryAspect
2 {
3     private TransactionScope scope;
4
5     public override void OnBefore(MethodArgs args)
6     {
7         this.scope = new TransactionScope(TransactionScopeOption.RequiresNew);
8     }
9
10    public override void OnSuccess(MethodArgs args)
11    {
12        this.scope.Complete();
13    }
14
15    public override void OnException(MethodArgs args)
16    {
17        Transaction.Current.Rollback();
18    }
19
20    public override void OnAfter(MethodArgs args)
21    {
22        this.scope.Dispose();
23    }
24 }
```

Listing A.10: RunInTransactionAspect

Then apply this aspect to the original code:

```
1 [RunInTransactionAspect]
2 public void SaveOrder(string customerName, string item1, string item2)
3 {
4     //same body of code to save..
5 }
```

Listing A.11: SaveOrder with Aspect

This will ensure the method will operate in the transaction safe manner. We can apply the reusable aspect to other operations that need to be transaction scoped.

A.4 Integrate With MS-Build System

Buffalo can be integrated with MS-Build, so weaving can be invoked automatically when a project is compiled from the Visual Studio IDE. Note that the following instructions are

just the bare minimum to get this working, a lot of bell and whistle are omitted.

MS-Build is integrated with Visual Studio IDE via configuration file. For example, a C# project has the associated .csproj, if open in a text editor you will see a line that references a different configuration file: Microsoft.CSharp.targets. This file in turn references Microsoft.Common.targets.

Each .NET version has its own Microsoft.Common.targets file. Depending on the version you are using, open up this file in a text editor. For example, for .NET 4.0, this file is located in C:\Windows\Microsoft.NET\Framework\v4.0.30319\Microsoft.Common.targets.

Under the Compile section, around line #2013, add the line to import the Buffalo.targets file as shown in figure A.1

```

1995 <!--
1996 =====
1997                                     Compile
1998 =====
1999 -->
2000 <PropertyGroup>
2001     <CompileDependsOn>
2002         ResolveReferences;
2003         ResolveKeySource;
2004         SetWin32ManifestProperties;
2005         _GenerateCompileInputs;
2006         BeforeCompile;
2007         _TimeStampBeforeCompile;
2008         CoreCompile;
2009         _TimeStampAfterCompile;
2010         AfterCompile
2011     </CompileDependsOn>
2012 </PropertyGroup>
2013 <Import Project="C:\Buffalo\Buffalo.targets" />

```

Figure A.1: Adding Buffalo.targets

If you open Buffalo.targets in a text editor, it contains the following context:

```

1 <Project xmlns="http://schemas.microsoft.com/developer/msbuild/2003">
2   <PropertyGroup>
3     <CompileDependsOn>
4       $(CompileDependsOn);
5       Buffalo
6     </CompileDependsOn>
7   </PropertyGroup>
8

```

```
9    <Target Name="Buffalo">
10      <Message Text="Hello Buffalo! @(IntermediateAssembly)"/>
11      <Exec Command="&quot;C:\Buffalo\BuffaloAOP.exe&quot; &quot;@(IntermediateAssembly)&quot;"/>
12    </Target>
13 </Project>
```

Listing A.12: Buffalo.targets

This is how Buffalo get hooked into MS-Build, what this mean is that when user compiles a project, everything defined in the CompileDependsOn property group will be performed first, then a new target named "Buffalo" will be called immediately, which will invoke the BuffaloAOP via the Exec Command. Note that for the Exec Command, a complete path to BuffaloAOP.exe must be provided, including the encoded quotation marks as shown.

Make sure to save all the changes.

Now every time a C# project is compiled, Buffalo will be invoked automatically to perform the weaving.