

Buffalo: An Aspect Oriented Programming Framework for C#

by

Wei Liao

A Project Report Submitted
in
Partial Fulfillment of the
Requirements for the Degree of
Master of Science
in
Computer Science

Supervised by

Prof. James E. Heliotis

Department of Computer Science

B. Thomas Golisano College of Computing and Information Sciences
Rochester Institute of Technology
Rochester, New York

April 2013

The project “Buffalo: An Aspect Oriented Programming Framework for C#” by Wei Liao has been examined and approved by the following Examination Committee:

Prof. James E. Heliotis
Professor
Project Committee Chair

Prof. Matthew Fluet
Assistant Professor

Prof. Fereydoun Kazemian
Associate Professor

Dedication

To Jackson and Evan

Acknowledgments

I am grateful for my adviser Prof. Heliotis, whose insightful advice, guidance and support from the beginning not only enabled me to complete the project on time, but also led me to a better understanding of the subject area.

I am also grateful for Prof. Fluet and Prof. Kazemian for their invaluable feed backs.

Last but not the least, I want to thank my wife Michelle, for all the support and encouragement during my years at school.

Abstract

Buffalo: An Aspect Oriented Programming Framework for C#

Wei Liao

Supervising Professor: Prof. James E. Heliotis

Aspect Oriented Programming (AOP) is a paradigm that allows programmers to isolate and separate crosscutting concerns from their programs. The concept has not been widely adopted by modern languages; support in tooling such as Integrated Development Environment (IDE), is also rare. In this project, I designed and implemented Buffalo, an AOP framework to provide this capability for the .NET platform.

Buffalo performs Common Intermediate Language instruction set modification according to the aspects written by the developer with the help of the Mono Cecil library. Buffalo is .NET attribute based, which means developers with existing .NET skills will have little or no learning curve to get started. Buffalo will help increase developer productivity in many areas, such as unhandled exception catching, tracing and logging, and other areas.

Contents

Dedication	iii
Acknowledgments	iv
Abstract	v
1 Introduction	1
1.1 Compiler Support	2
1.2 Run-time Interception	2
1.3 Post Compilation Weaving	4
2 Design	5
2.1 A Buffalo Example	5
2.2 Usage Comparison With AspectJ and PostSharp	6
2.3 Aspect Interface	8
2.4 MethodArgs	9
2.5 A Buffalo Aspect	9
2.6 MethodBoundaryAspect	10
2.7 MethodAroundAspect	11
2.8 How to Apply an Aspect	13
2.9 Visual Studio Solution Structure	15
3 Implementation	17
3.1 Implementation Overview	17
3.2 MethodBoundaryAspect Implementation Detail	17
3.3 MethodAroundAspect Implementation Detail	19
3.4 MethodArgs Implementation Detail	21
4 Analysis	22

5	Conclusions	25
5.1	Current Status	25
5.2	Future Work	26
5.3	Lessons Learned	26
	Bibliography	27
A	User Manual	29
A.1	Compiling	29
A.2	Simple Profiler	30
A.3	Transaction Database	32
A.4	Update UI Safely with Thread	35
A.5	Integrate With MS-Build System	37

List of Figures

1.1	AOP Framework Using Proxy Pattern	3
1.2	Buffalo Model	4
2.1	Aspect Inheritance	8
2.2	Illustration of MethodAroundAspect	12
2.3	Diagram for Finding Eligible Methods	14
2.4	Solution Structure	16
3.1	CIL Interception Points	19
A.1	UI Thread Example	35
A.2	Adding Buffalo.targets	38

Chapter 1

Introduction

Object Oriented Programming (OOP) languages have given programmers much freedom in expressing themselves in Object Oriented Design. However, they are still lacking in some areas when it comes to particular software design decisions such as cross-cutting concern.

Cross-cutting concerns are aspects of the program that scatter throughout different areas of the code base. These aspects cannot usually be separated easily or cleanly from the rest of the program. The essence of Aspect Oriented Programming [1] is to find ways to solve the cross-cutting concern problems.

In this project, a framework called "Buffalo" is designed and implemented to solve this type of problem on the .NET platform. Buffalo makes use of the .NET attribute system to weave aspect code to any targeted methods. The design and rationale of the framework is discussed in Chapter 2. The implementation detail is given in Chapter 3.

The results indicate that by using Buffalo, developers can separate cross-cutting concerns from the core of the program for easy maintenance and ultimately be more productive. The analysis is discussed in Chapter 4.

The report concludes in Chapter 5 with the current project status. A set of planned future works is also discussed as well as what was learned from working on this project.

Buffalo comprises around 1,200 lines of source code. A user manual is included in Appendix A, which contains examples on how the system works. Instructions are also included for how to integrate Buffalo with MS-Build.

1.1 Compiler Support

There are several broad approaches to implementing an AOP framework. The ideal approach is to extend the compiler of the target language to provide built-in support, thus making AOP the first class citizen. However there are very few languages out there that take this approach; among the few are Delphi Prism [2] and AspectJ [3, 4].

Microsoft is currently in the "wait and see" mode regarding support of AOP development in the C# compiler [5]. The alternative compiler, such as Mono C# [6], is open source, so technically, anyone can build AOP support into it. While that would have been a fun challenge, it would have been a fairly big undertaking, and there is concern that the project might not be finished in the time frame required.

That leaves framework support as the other viable option. There are several implementation techniques to provide AOP capabilities [7, 8, 9] via framework.

1.2 Run-time Interception

Early on the implementation, approaches were narrowed down to two approaches: Run-time Interception and Post Compilation Weaving. As its name suggested, Run-time Interception operates while the program is in execution. It uses the proxy pattern where client communicates with the target object via a proxy, and aspects are injected to the proxy. This enables run-time behavior of the program to be modified. Figure 1.1 illustrates this process.




Figure 1.1: AOP Framework Using Proxy Pattern

New functionality can be added to the target object via the proxy. The disadvantage of this approach is that it involves the generation of proxy object at run-time. As a result, the run-time performance of the application will be impacted. It is also restricting in that both target object and the proxy must implement a common interface for this to work and that only virtual methods are exposed for interception.

From the end user's perspective, to use it, the developer usually must provide some type of mapping between the target object and the proxy via a configuration file so that the actual proxy generation can occur.

Although this approach is easier to implement, it is not as easy and user-friendly. Buf-falo is not taking this approach mainly because one of the goals is to keep it flexible and simple to use.

1.3 Post Compilation Weaving

The approach Buffalo takes is Post Compilation Weaving. The idea is that after compilation of the source code, Buffalo takes over to disassemble the assembly and weaves in the defined aspect code to all targeted methods. This approach is more difficult to implement as it involves modifying the underlying assembly by changing Common Intermediate Language (CIL) instructions [10]. However, the advantage is that no run-time performance of proxy generation will be needed. Therefore, no messy configuration is required.

Since injection happens post-compilation, the whole process can be integrated into the MS-Build system to have the weaving invoked automatically if needed. This will further reduce the steps needed from the developer.

Figure 1.2 shows an overview of the compilation process and where Buffalo will fit in the process.

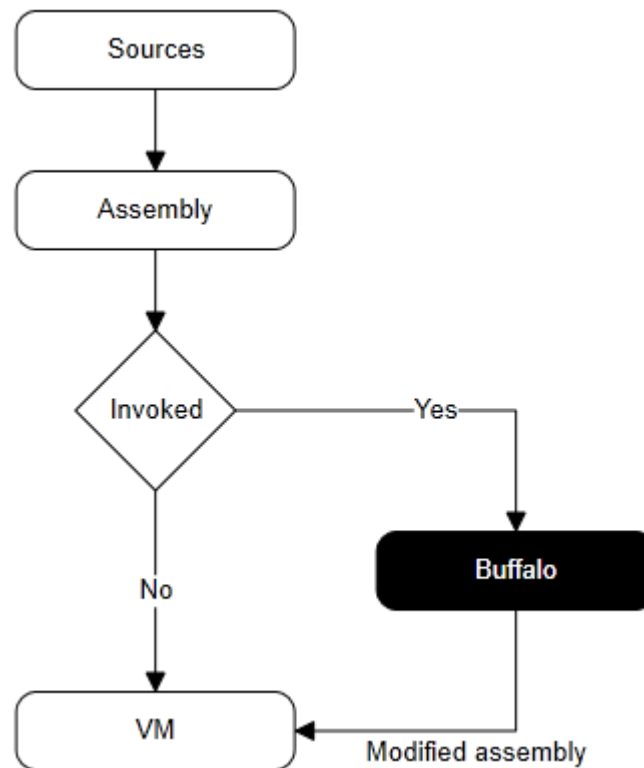


Figure 1.2: Buffalo Model

Chapter 2

Design

2.1 A Buffalo Example

An example of how to use Buffalo is presented in Listing 2.1. Further detail on the design will be discussed in later sections. The example shows how a simple aspect is created. It inherits `MethodBoundaryAspect`, which is a predefined aspect by Buffalo, and overrides its `OnBefore` and `OnAfter` functions.

```
1 using Buffalo;
2 using System;
3 public class TraceAspect : MethodBoundaryAspect
4 {
5     public override void OnBefore(MethodArgs args)
6     {
7         Display("ENTERING", args);
8     }
9     public override void OnAfter(MethodArgs args)
10    {
11        Display("EXITING", args);
12    }
13    void Display(string title, MethodArgs args)
14    {
15        Console.WriteLine("{0} {1}", title, args.FullName);
16        foreach (var p in args.Parameters)
17        {
18            Console.WriteLine("\t{0} ({1}) = {2}",
19                               p.Name, p.Type, p.Value);
20        }
21    }
22 }
```

Listing 2.1: Sample TraceAspect

To use the aspect, simply apply it to a method, class or assembly. Once the code is compiled to produce an assembly, BuffaloAOP.exe can be invoked by passing it the path to the assembly. Buffalo will take over and weave in the aspect. This and more examples and details are also provided in Appendix A.

```
1 [TraceAspect]
2 public class Hello
3 {
4     //...
5 }
```

Listing 2.2: Apply Aspect on Class Level

This ensures that whenever a public function is invoked by an instance of Hello, OnBefore and OnAfter of the aspect will be invoked respectively.

2.2 Usage Comparison With AspectJ and PostSharp

Buffalo is influenced by AspectJ and PostSharp. PostSharp makes use of System.Attribute as well, so creation and usage of aspect is similar to Buffalo. A simple example to replace a function call foo() is presented below to illustrate some of the similarities and differences.

Listing 2.3 shows that Buffalo creates an aspect inheriting from MethodAroundAspect, which is one of the predefined aspect in Buffalo. It overrides the Invoke(), so anytime the annotated function is invoked, the aspect will be called to return 3 instead of 100.

```
1 public class MyAspect : MethodAroundAspect
2 {
3     public override object Invoke(Buffalo.MethodArgs args)
4     {
5         return 3;
6     }
7 }
8
9 ...
```

```
10 public class Program
11 {
12     [MyAspect]
13     public int foo()
14     {
15         return 100;
16     }
17 }
```

Listing 2.3: Buffalo MethodAroundAspect

AspectJ is a compiler itself that produces compatible Java bytecode, and it has its own syntax that is somewhat different from Java. Listing 2.4 defines an aspect that will replace all the calls to `foo()` and returns the integer 3 instead.

```
1 aspect A {
2     int around(): call(int foo()) {
3         return 3;
4     }
5 }
```

Listing 2.4: AspectJ Around

Listing 2.5 shows an example of PostSharp, which is very similar to how Buffalo is used. The big difference is that the return value is carried in the argument object instead of returning it directly.

```
1 [Serializable]
2 public class MyAspect : MethodInterceptionAspect
3 {
4     public override void OnInvoke(MethodInterceptionArgs args)
5     {
6         args.ReturnValue = 3;
7     }
8 }
9
10 public class Program
11 {
12     [MyAspect]
13     public int foo()
14     {
```

```

15         return 100;
16     }
17 }

```

Listing 2.5: PostSharp MethodInterceptionAspect

2.3 Aspect Interface

Figure 2.1 shows the relationship of various aspect types in Buffalo. This is used by Buffalo to identify aspects during reflection.

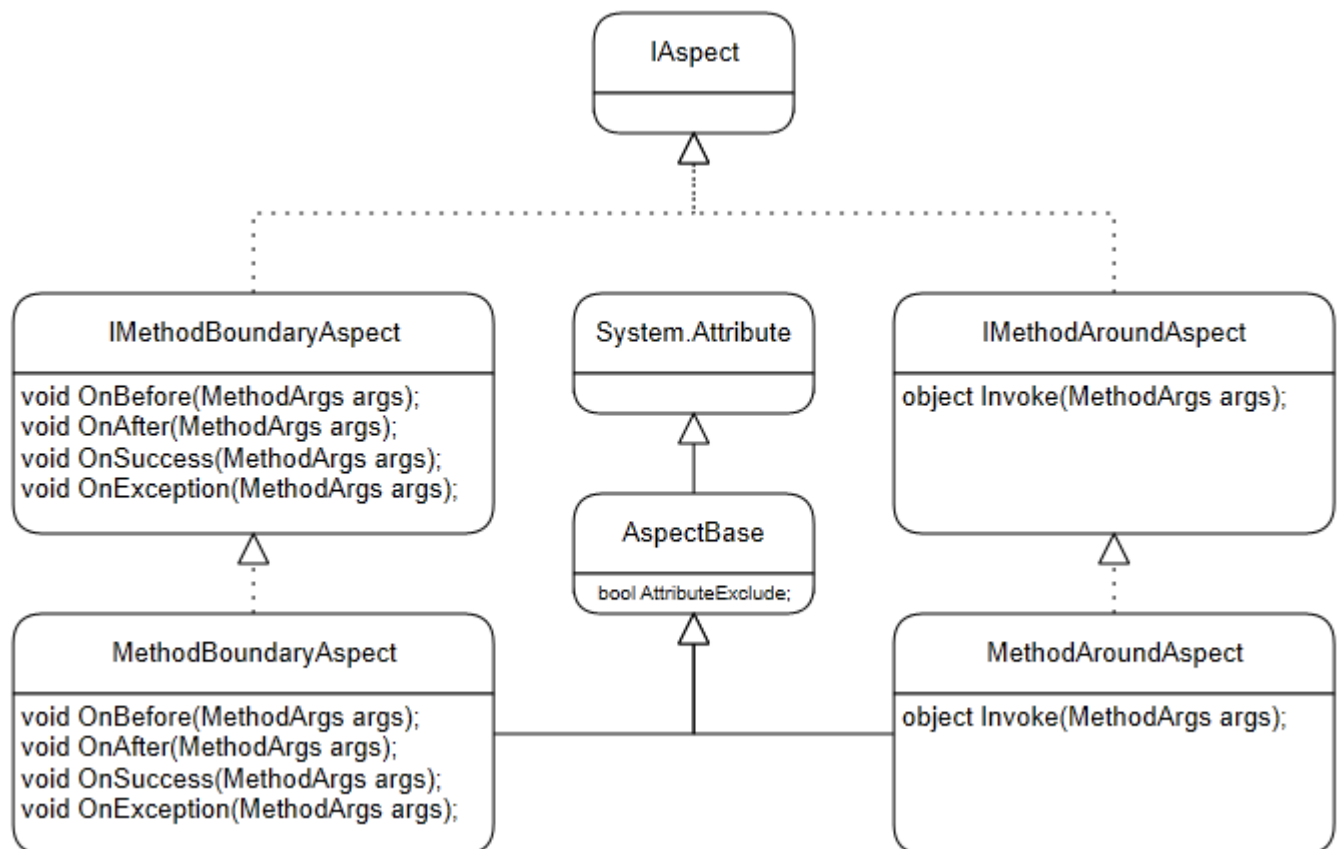


Figure 2.1: Aspect Inheritance

All aspects ultimately implement the IAspect interface; therefore, it can be reasoned that for all the public types in an assembly that if a type implements IAspect, then it must be an aspect itself.

Buffalo supports more than one aspect applied at any given level. This will allow developers more flexibility while developing multiple aspects and applying them as needed.

Furthermore, by default, an aspect will be automatically excluded from applying to itself. This is implemented to prevent stack overflow in some cases. Although argument can be made that an aspect should be able to be applied to a different aspect; that is not currently implemented in Buffalo.

2.4 MethodArgs

As mentioned above, when all is said and done, an aspect ultimately gets injected into each *individual* method. When developing an aspect, meta-information about the target method can be accessed. This information is encapsulated via the MethodArgs object passed in as parameter to the aspect. Method name, its full method signature, return type and parameter list (including parameter name, type and value) are captured for each target method.

The parameter list capturing is especially of interest; it enables the developer to peek inside the method that is executing at various point and inspect its parameter values. This will be useful in cases such as exception handling, where it will be useful to see what the actual values were at the time of the exception. When using MethodAroundAspect, the parameter values can even be modified and passed back to the original method.

2.5 A Buffalo Aspect

When performing Post Compilation Weaving, Buffalo must have the ability to discover what aspect is applied to what methods in an assembly. In order to achieve that, the target assembly has to carry some identifying meta-data.

A given .NET assembly already carries a great deal of such meta-data for various purposes. .NET has the System.Attribute type that exists primarily for the purpose of inserting meta-data into the assembly during compilation. When the source code is compiled, it is converted into CIL [11] and put inside a portable executable (PE) file, with the meta-data

generated by the compiler.

Buffalo takes advantage of this characteristic in two phases.

1. An aspect defined in Buffalo will be in the form of an attribute by inheriting from `System.Attribute`. It can, therefore contain any valid .NET code. Specifically, however, an aspect needs to override various predefined methods in order to do something useful. The next section will discuss the relationship between various aspect types.
2. After compilation, the assembly will now contain the meta-data about all the aspects. Buffalo can inspect the assembly for such information and perform CIL code injection accordingly.

In other words, a Buffalo aspect is a .NET attribute in disguise.

2.6 MethodBoundaryAspect

What functionality does Buffalo support? What type of weaving does it do? For inspiration, existing works, such as AspectJ [3] and PostSharp [8], were studied. Specifically, Buffalo will intercept the various points of an executing method. Those points are namely the following: before a method executes; after a method executes; whether or not the method executed successfully without error; or whether the method throws an exception at any point during the execution. These various points of interception are grouped into the `MethodBoundaryAspect`.

`MethodBoundaryAspect` can be cleanly mapped to the try-catch-finally statements of the .NET languages. As far as the runtime is concerned [12, 13], try-catch can be used liberally without serious performance degradation. For example, a simple method is shown in Figure 2.6:

```
1 public void SomeFunction () {  
2     //Perform some action...  
3 }
```

Listing 2.6: Sample Function

When performing CIL modification, the above will be transformed by Buffalo into something shown in Figure 2.7. This clearly captures the spirit of the `MethodBoundaryAspect`. Regardless of whether the source already contains its own try-catch, or try-catch-finally blocks, Buffalo will wrap the body of the method inside of a new try-catch-finally block.

```
1 public void SomeFunction () {  
2     try {  
3         OnBefore();  
4         //Perform some action  
5         OnSuccess();  
6     }  
7     catch (Exception e) {  
8         OnException(e);  
9     }  
10    finally {  
11        OnAfter();  
12    }  
13 }
```

Listing 2.7: Sample Try-catch-finally

The transformed method in Figure 2.7 still does what the original method intends to do; only now at various points, execution is being intercepted to provide more functionality.

2.7 MethodAroundAspect

Another type of aspect that Buffalo supports is the `MethodAroundAspect`. Rather than intercepting various execution points of a method, the method can be completely replaced by another method defined in an aspect while preserving the option to call back into the original method if necessary.

At first glance `MethodAroundAspect` sounds straightforward to implement, but it turns out to be much more involved than the `MethodBoundaryAspect`.

Since the option to call back into the original method is needed, it is critical that the original method is not modified under any circumstance. If the method body instructions

are simply overridden with that of the replacement, the call back to the original method will be meaningless since the method is now changed. The original method must stay intact during the CIL modification.

To get around this obstacle, whenever Buffalo encounters the `MethodAroundAspect` applied to a method, it dynamically generates a replacement method in CIL with the same method signature as the original.

The body of this replacement method is also completely different than the original. It instantiates the aspect and makes a call to the `Invoke()` method, which is the actual code that will be ran as a replacement.

Inside the `Invoke` method, the developer can issue a call back to the original method via a call to the `Proceed()` method. Then throughout the program, for any calls made to the original method, Buffalo will change them to call the replacement method instead. This process is illustrated in Figure 2.2.

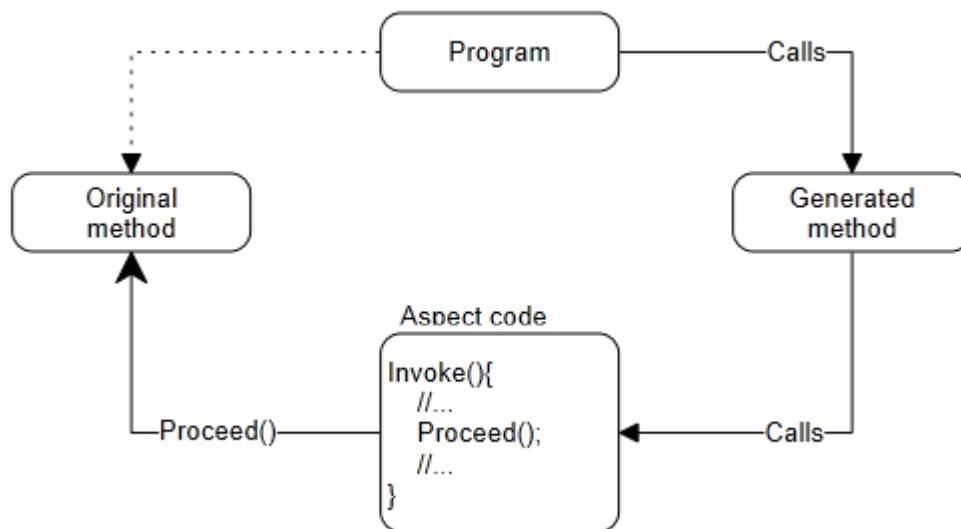


Figure 2.2: Illustration of `MethodAroundAspect`

The dotted line from `Program` to the original method indicates that once the `MethodAroundAspect` is applied to it, from the perspective of CIL the program cannot directly access that method any more. Access to the original method now must come from inside

the aspect. Also note that the original method is not changed at any given time.

2.8 How to Apply an Aspect

Since an aspect is really a .NET attribute, it can be used just like any other attribute. However, code annotated with an aspect is special in that it can be understood only by Buffalo.

A Buffalo aspect can be applied on three levels, with the following characteristics:

1. Method - apply the aspect to an individual method.
2. Class - if aspect is applied to a class, all public methods, including the public properties, are automatically applied.
3. Assembly - if aspect is applied to an assembly, #2 will apply for all of the public classes within the assembly.

An exception to the above rule is the `MethodAroundAspect`, where it can only be applied on a method level, as will be shown later.

All aspects have a property named `AttributeExclude`; if this property is set to true, then the annotated target will not be included in the weaving. This exclusion can happen on any level. For example, if a method contains this annotation `[SampleAspect(AttributeExclude=true)]`, the method will be skipped for the `SampleAspect` during the weaving process.

No matter how the aspect is applied, ultimately it will result in a list of the methods that are annotated. This simply means that if the aspect is applied to a single method, that method is the only one that will get CIL modified. If the aspect is applied on the whole assembly, then all public methods will be CIL modified.

To get the list of the eligible methods for CIL modification, Buffalo attempts various checking according to Figure 2.3 to see if it should include a given method.

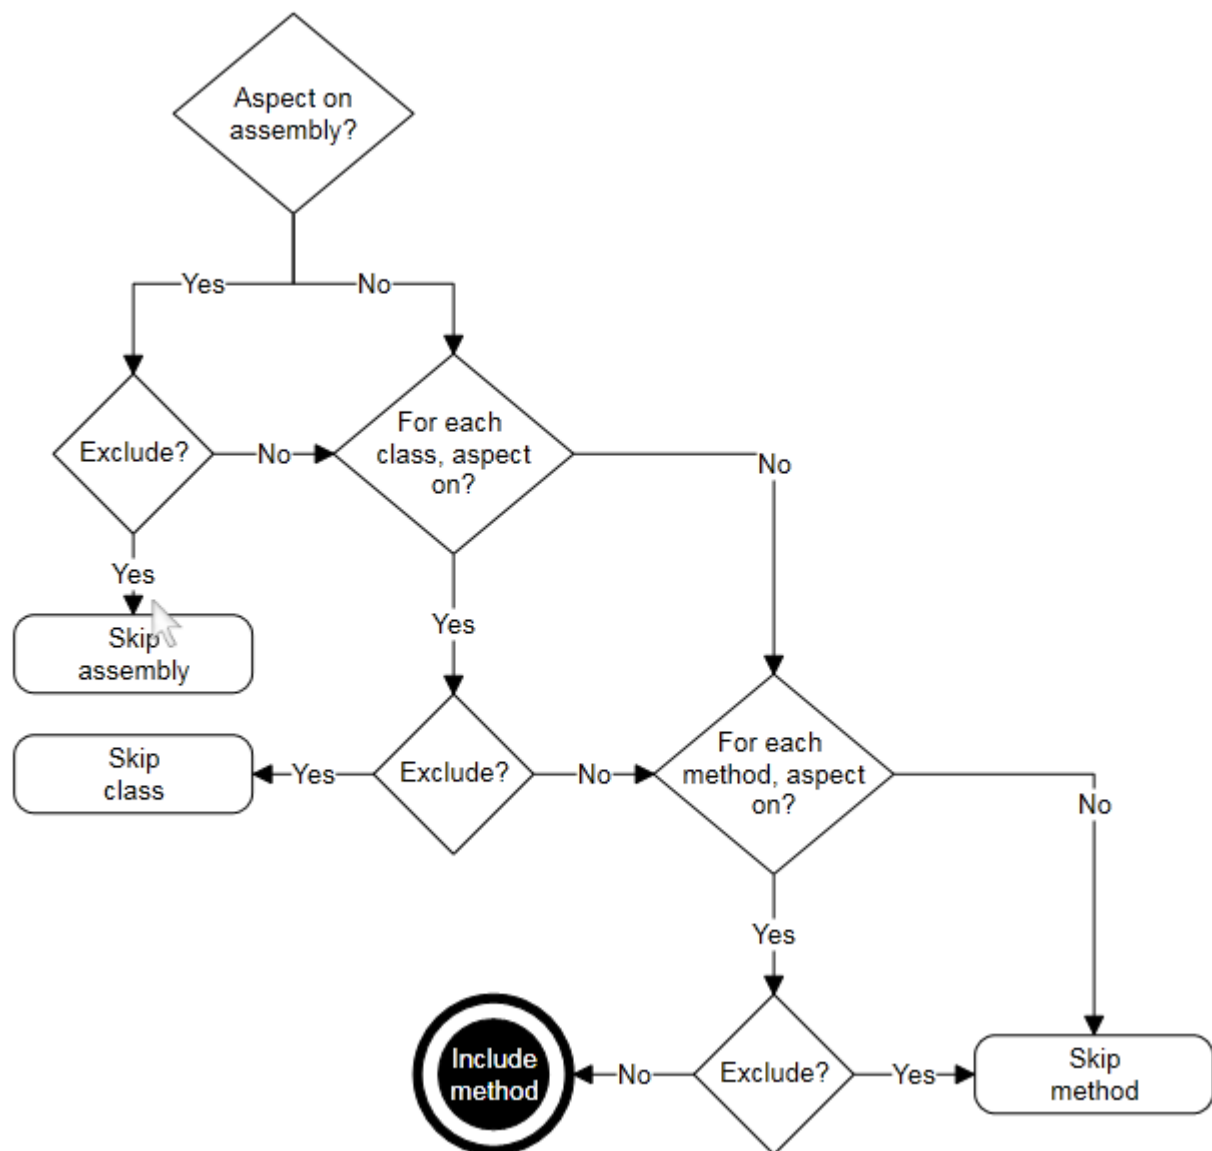


Figure 2.3: Diagram for Finding Eligible Methods

If no aspect is applied on the assembly, it does not necessarily mean no aspect is applied anywhere; the aspect might still be applied on any given class or method.

Buffalo first checks if an aspect is applied to the target, and then checks if it is set to be excluded. At the end it will end up with a list of methods that should be CIL modified.

2.9 Visual Studio Solution Structure

Originally Buffalo was implemented as one executable; that includes the various aspects and the program that initiates the weaving. It was later separated into two assemblies. One is a class library that contains the actual implementation. Another is a command line executable that calls into the class library to perform the weaving. This separation is necessary so that the developer can perform weaving from the command line or hook into MS-Build if necessary.

To actually write the aspect, one only needs to reference the class library as underlined in Figure 2.4.

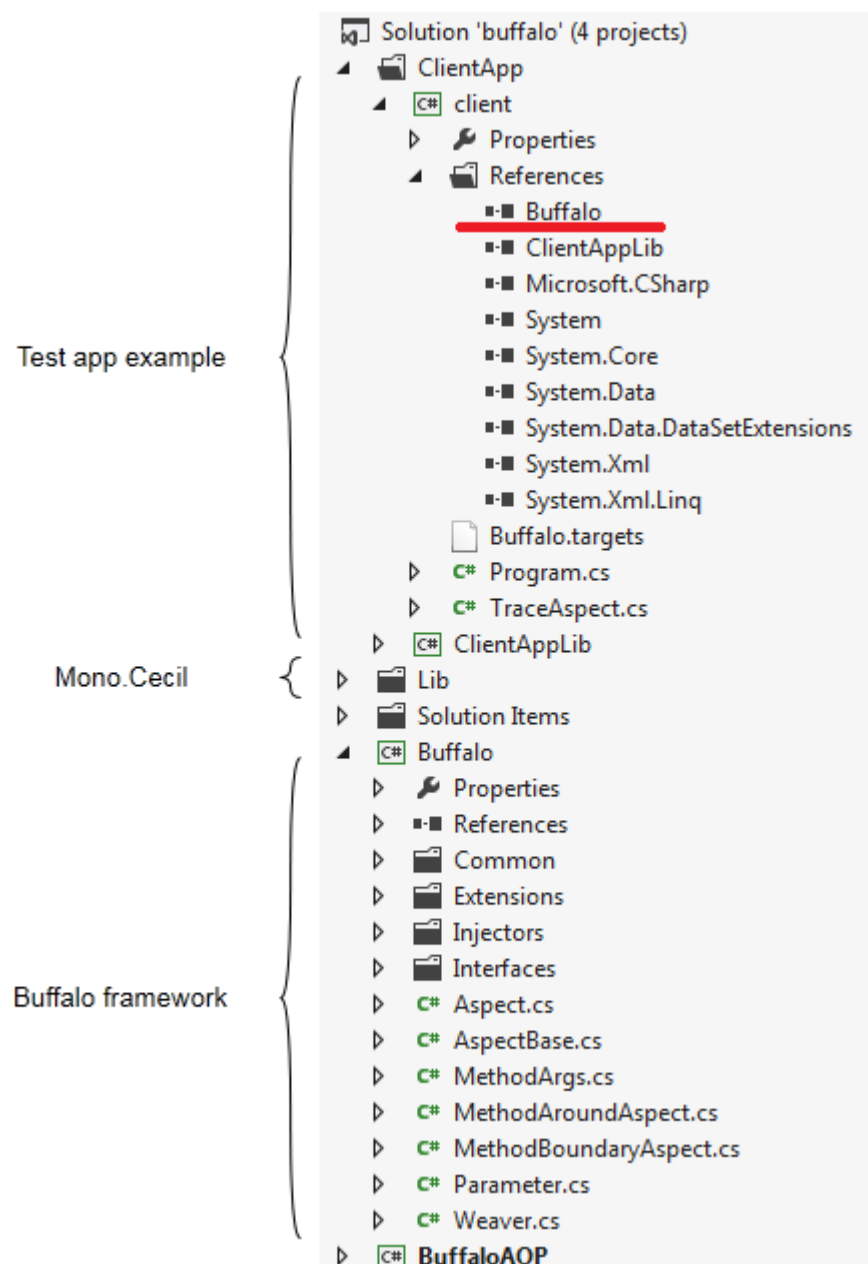


Figure 2.4: Solution Structure

The client project shown above is a simple program included in the solution for testing.

Chapter 3

Implementation

3.1 Implementation Overview

The implementation begins by finding all eligible methods to be injected using the verification process indicated in Figure 2.3. Each eligible method will have one or more aspects applied to it. The actual injection process will loop through each aspect for the eligible method and inject the necessary CIL instructions.

The CIL instructions modification is performed by using the open-source Mono.Cecil library. Originally, the System.Reflection APIs in the .NET Framework were favored, but it was later determined that Mono.Cecil provides more features and is much easier to work with. Buffalo uses Mono.Cecil heavily to modify CIL instructions and to assemble the final assembly.

3.2 MethodBoundaryAspect Implementation Detail

Each type of aspect has its own injector that implements the `IInjectable` interface. This interface contains only one method contract - `Inject(..)`. It takes the list of eligible methods and injects the appropriate aspect to them.

`MethodBoundaryAspect` is pretty straightforward to implement. Take the following hello-world example wrapped in a try-catch-finally block as mentioned previously:

```
1 public void SayHello()  
2 {  
3     try{  
4         Console.WriteLine(Hello World!);
```

```
5     }catch(Exception ex){  
6     }finally{  
7     }  
8 }
```

Listing 3.1: SayHello Function

The generated CIL is shown in Figure 3.2. For ease of display, the CIL has been cleaned up a bit:

```
1  .try  
2  {  
3      .try  
4      {  
5          IL_0002: Ldstr "Hello World!"  
6          IL_0007: call void [mscorlib]System.Console::WriteLine(string)  
7          IL_000e: leave.s IL0015  
8      }  
9      catch [mscorlib]System.Exception  
10     {  
11         IL_0010: stloc.0  
12         IL_0013: leave.s IL_0015  
13     }  
14     IL_0015: leave.s IL_001c  
15 }  
16 finally  
17 {  
18     IL_001a: endfinally  
19 }  
20 IL_001c: ret
```

Listing 3.2: CIL Generated for Sample C# Function

Figure 3.2 shows the standard emission of the CLR when it encounters the try-catch-finally statement. In CLR, there is a concept of the protected region, where each region is associated with a handler. A try-catch-finally is actually encapsulated in two such regions: a catch and a finally. From here, it can be easily figured out where to inject the various boundary aspects, as shown in Figure 3.1.




Figure 3.1: CIL Interception Points

3.3 MethodAroundAspect Implementation Detail

The MethodAroundAspect, on the other hand, is a much more complicated when compared to the MethodBoundaryAspect.

MethodAroundAspect implements IMethodAroundAspect which has the following contract:

```

1 internal interface IMethodAroundAspect : IAspect
2 {
3     void Invoke(MethodArgs args)
4 }

```

Listing 3.3: IMethodAroundAspect Interface

In the development of an aspect, a Proceed() can be issued to signal a call back into the original method. The steps taken to implement MethodAroundAspect in CIL are roughly

as follows:

1. Create a replacement for the annotated function with exactly the same method signature.
2. Create and store a variable pointing to the aspect.
3. Copy all parameters from original method to the newly created replacement function.
4. Create a variable to hold MethodArgs.
5. Issue a call to Invoke() from the replacement function, passing in the MethodArgs variable.
6. Handle the return value appropriately.
 - (a) If the original method returns non-void type, then put the return value back on the stack.
 - (b) If the original method returns void, discard the return value from Invoke().
7. Handle Proceed() that might be issued from inside the Invoke().
 - (a) Load all the parameters onto the stack.
 - (b) Call back into the original method.
 - (c) Handle the return value appropriately.
8. Modify all calls from original method to the replacement method.

As Figure 2.2 shows, the actual calling of either the original or replacement method is abstracted away. This is also a testament of the saying in software engineering that "anything can be resolved by another layer of abstraction."

Another important distinction is that MethodAroundAspect currently can be applied only on the method level, and that it should be applied to one method only. This is by design because a replacement method might not be appropriate to replace more than one

method. Especially if the aspect is applied on the assembly level, all the public methods will be replaced by a single replacement method.

3.4 MethodArgs Implementation Detail

When an aspect is being developed, information about the target method can be accessed. This is achieved by using the MethodArgs object. During the weaving, an instance of MethodArgs is created with all properties assembled dynamically to capture the information of the current executing method. MethodArgs is then passed as a parameter into each of the aspect.

Being able to capture some information about the annotated methods will be extremely useful in cases where debugging becomes necessary.

A distinct instance of MethodArgs for each boundary aspect was instantiated in an early Buffalo implementation. Later on as an optimization, only one instance is instantiated at the beginning of the method body, and that instance is used in all the boundary aspects for a target method.

An example of how to use MethodArgs is presented in the user manual.

Chapter 4

Analysis

The project hypothesized that by using Buffalo, programmers can separate the cross-cutting concerns from their applications quickly and easily. While the current implementation of Buffalo cannot be shoehorned into every situation, in many scenarios cross-cutting concerns can be decoupled and encapsulated into distinct units of code as an aspect, represented as a .NET attribute type. This separation of concern is what really encourages the code re-usability. Buffalo facilitates this process by providing the necessary plumbing.

While it is difficult to quantify the benefit of using Buffalo, it offers as a solid alternative to managing the cross-cutting problem. It enables aspect creations to capture and isolate problems that are spreading into different modules of the program. Because Buffalo makes use of the System.Attribute type that already existed on the .NET platform, developers will have the full power of the .NET framework at their disposal when creating the aspects.

Getting started with Buffalo is also relatively easy compared to other frameworks that are configuration-based, where an interface must map to a concrete implementation in order for code injection to work. Buffalo has no such configuration-file requirement. Buffalo is 100% managed code; developers already familiar with the .NET framework will have a minimal learning curve.

Although the title of this report indicates that the framework is developed in C#, by no means is it limited to the C# language. Since Buffalo works with the Common Intermediate Language, in practice Buffalo can be used with other .NET languages that compile to CIL, such as VB.NET. Preliminary tests show that Buffalo can also work in the Linux environment where the Mono C# compiler is installed.

Generally speaking, using Buffalo it also makes the code a bit cleaner to read in that it abstracts away the common code into aspects. Since aspects are centralized, maintainability is much easier.

Buffalo can be useful in situations where cross-cutting concerns are present. One common case is exception logging, where an aspect can be developed to catch all the unhandled exceptions.

One can argue that since unhandled exceptions will bubble up the chain, a developer can simply catch them in the main method, and this would have achieved a similar effect. However, that approach is limited in that it is very generic. When the main method catches an exception, it has no idea what the internal state of the failed method was. Using Buffalo, the internal state can be inspected by checking the `MethodArgs` object, giving the developer a better sense at what caused the failure.

Besides exception handling, another common problem facing developers working on databases is to ensure data atomicity. Imagine a customer who deposits money to a bank. The amount is saved successfully to the bank's own record, but the customer's account failed to get updated due to an unforeseeable network error. Not only would the database end up with incomplete data, but the bank would have a very unhappy customer. Either all operations are successfully completed or nothing should have been written to the database.

Buffalo can be used to solve this problem by encapsulating transactional scope using the `MethodBoundaryAspect`. A complete example is available in Appendix A.

Developers that work with UI elements often run into problems updating the UI from a different thread. As UI elements can be updated only by the thread that created them, this creates the cross-threading problem. Buffalo's `MethodAroundAspect` can be used to get around this problem by marshaling the update back to the UI thread. For more detail please refer to Appendix A.

Buffalo allows developers to create aspects quickly to solve various problems. Besides the scenarios mentioned above, it can also be used for authorization, where security credentials are validated before method execution. It can be used to create a caching mechanism

to improve application performance. Buffalo has performed well in isolating cross-cutting concerns into single units of code, which are easily maintained and modified.

Chapter 5

Conclusions

5.1 Current Status

Buffalo currently contains two types of aspect: `MethodBoundaryAspect` - where various execution points can be intercepted - and `MethodAroundAspect` - where a method can be wrapped around or completely replaced while preserving the option to call back into the original method. `MethodAroundAspect` targets the individual method, whereas `MethodBoundaryAspect` can be applied on three levels on a .NET application.

Originally MS-Build integration was planned. When the Buffalo program was installed via the setup executable, it was intended for it to modify the relevant .NET configuration files to hook `BuffaloAOP.exe` into MS-Build. That would trigger the weaving process from within the Visual Studio IDE when the solution is compiled. It was later found that Microsoft has dropped support for the setup project type. As a work around, instruction on how to hook into MS-Build manually is provided in Appendix A.

Originally, `MethodAroundAspect` was intended to be used similar to `MethodBoundaryAspect`, where it could be applied on three levels. However, it was later determined that since the CIL instruction of the actual aspect will also have to be modified, it really does not make sense for it to be applied to more than one method as it would introduce conflicting changes.

`MethodAroundAspect` improvement remains to be determined.

5.2 Future Work

There are several areas where Buffalo can be improved. Usability-wise, there is currently no automatic setup program that installs Buffalo onto users' computer. There is also no automatic integration into MS-Build System. Some manual steps are still needed in order to provide a more seamless experience.

Functionality-wise, when Buffalo performs Post Compilation Weaving it starts fresh each time; it would be interesting to see how incremental weaving can be done here. Another useful functionality is to be able to apply aspects by matching a set of methods.

5.3 Lessons Learned

A framework such as Buffalo mitigates the problem of cross-cutting concerns. Still, to tackle the root of the problem efficiently, compiler vendors must actively embrace the AOP concept and provide native support in their programming languages.

Developers also have to understand such problems and what solutions are available to better educate themselves.

Only when the concept is widely understood and supported by both developers and vendors can there hope to begin alleviating such problems.

Bibliography

- [1] et al. Gregor Kiczales, John Lamping. Aspect-Oriented Programming. In *Proc. of European Conference on Object-Oriented Programming (ECOOP)*, Finland, 1997.
- [2] Cirrus - The Oxygene Language Wiki. <http://wiki.oxygenelanguage.com/en/Cirrus>. Accessed: 07/16/2012.
- [3] The Eclipse Foundation. AspectJ Frequently Asked Questions. <http://www.eclipse.org/aspectj/doc/released/faq.php>. Accessed: 07/19/2012.
- [4] Ramnivas Ladda. *Aspectj in Action: Enterprise AOP with Spring Applications*. Manning Publications, second edition, October 2009.
- [5] MSDN TV. Transcript: Whiteboard with Anders Hejlsberg. <http://msdn.microsoft.com/en-us/cc320411.aspx>. Accessed: 07/19/2012.
- [6] Mono. CSharp Compiler - Mono. http://www.mono-project.com/CSharp_Compiler. Accessed: 07/19/2012.
- [7] M. Devi Prasad and B. D. Chaudhary. AOP Support for C#. In *Proc. of Second International Conference on Aspect-Oriented Software Development*, pages 49–53, 2003.
- [8] Gael Fraitteur. User-friendly aspects with compile-time imperative semantics in .NET: an overview of PostSharp. In *International Conference on Aspect-Oriented Software Development*, 2008.
- [9] Howard Kim. AspectC#: An AOSD implementation for C#. Master’s thesis, Trinity College Dublin, 2002.
- [10] Aleksandr Mikunov. Rewrite MSIL Code on the Fly with the .NET Framework Proling API. *MSDN Magazine*, 2003.
- [11] Serge Lidin. *Expert .NET 2.0 IL Assembler*. Apress, first edition, August 2006.

- [12] ECMA International. ECMA-334 - C# Language Specification. <http://www.ecma-international.org/publications/files/ECMA-ST/Ecma-334.pdf>, June 2006.
- [13] ECMA International. ECMA-335 - Common Language Infrastructure (CLI). <http://www.ecma-international.org/publications/files/ECMA-ST/Ecma-335.pdf>, June 2012.

Appendix A

User Manual

It is easy to get start using Buffalo. I assume you are reasonably familiar with the Visual Studio IDE and the general layout of a VS solution. I also assume you know how to compile a solution and where to find the compiled assembly. Buffalo is developed in VS2012; it is recommended that you have the same version of the IDE installed. The instructions and examples mentioned in this manual are specifically for the Windows platform. Although preliminary tests show that the project can be compiled under Linux when Mono is installed.

A.1 Compiling

To begin, first download the full source code from <https://github.com/wliao008/buffalo>. Open buffalo.sln in VS2012, and compile the source code. This will produce the Buffalo.dll and BuffaloAOP.exe in their respective bin/debug folders.

For this example, we will perform the weaving from a command prompt. Create a folder name under the C drive. Here, I will call it "Buffalo." Copy Buffalo.dll, BuffaloAOP.exe and Mono.Cecil.dll to C:\Buffalo. Note that this folder can be located anywhere in your system; I am just putting it on the C drive for simplicity.

Now let us create an aspect.

A.2 Simple Profiler

In this example we will create a profiler for our application. Suppose we have the following simple program:

```
1 using System;
2
3 namespace Hello
4 {
5     class Program
6     {
7         static void Main(string[] args)
8         {
9             Hello h = new Hello();
10            h.SayHello();
11            h.Say("Hey Buffalo how's it going!");
12
13            //pause the console
14            Console.Read();
15        }
16    }
17
18    public class Hello
19    {
20        public void SayHello()
21        {
22            Console.WriteLine("Hello World!");
23        }
24
25        public void Say(string msg)
26        {
27            Console.WriteLine(msg);
28        }
29    }
30 }
```

Listing A.1: Hello Program

When the program runs, it will display the following output:

```
1 Hello World!
2 Hey Buffalo how's it going!
```

Listing A.2: Hello Program Output

Also suppose that we want to monitor the program. We want to know when a method was accessed and exited. We can easily create an aspect to do such work:

```
1 using Buffalo;
2 using System;
3
4 public class TraceAspect : MethodBoundaryAspect
5 {
6     public override void OnBefore(MethodArgs args)
```

```

7      {
8          Display("ENTERING", args);
9      }
10
11     public override void OnAfter(MethodArgs args)
12     {
13         Display("EXITING", args);
14     }
15
16     public override void OnSuccess(MethodArgs args)
17     {
18         Display("SUCCESSFULLY EXECUTED", args);
19     }
20
21     public override void OnException(MethodArgs args)
22     {
23         Display("EXCEPTION ON", args);
24     }
25
26     void Display(string title, MethodArgs args)
27     {
28         Console.WriteLine("{0} {1}", title, args.FullName);
29         foreach (var p in args.Parameters)
30         {
31             Console.WriteLine("\t{0} ({1}) = {2}", p.Name, p.Type, p.Value);
32         }
33     }
34 }

```

Listing A.3: TraceAspect

With the aspect defined, now we can apply this aspect on any of the three different levels. Let's apply it to the Hello class for example.

```

1 [TraceAspect]
2 public class Hello
3 {
4     //...
5 }

```

Listing A.4: Apply Aspect to the Hello Class

Now everything is in place. We can now invoke the BuffaloAOP.exe to perform the weaving. Open a command prompt, and navigate to C:\Buffalo. Issue this command:

```

1 C:\Buffalo>BuffaloAOP.exe <path.to.the.hello.program.exe>

```

Listing A.5: Invoking BuffaloAOP.exe

Replace path to the hello program exe with the actual complete path to the program assembly. Suppose the program assembly is located at C:\Projects\Hello\bin\Hello.exe; we would issue the command as follows:

```
1 C:\Buffalo>BuffaloAOP.exe C:\Projects\Hello\bin\Hello.exe
```

Listing A.6: Invoking BuffaloAOP.exe Example

If everything goes well BuffaloAOP.exe will perform the injection and put the final assembly in the Modified folder inside the folder of the target assembly. In this case, it will be at C:\Projects\Hello\bin\Modified\Hello.exe. Now when the program runs, it will display the following output:

```
1 ENTERING System.Void Hello.Program::Main(System.String[])
2     args (System.String[]) = System.String[]
3 ENTERING System.Void Hello.Hello::.ctor()
4 SUCCESSFULLY EXECUTED System.Void Hello.Hello::.ctor()
5 EXITING System.Void Hello.Hello::.ctor()
6 ENTERING System.Void Hello.Hello::SayHello()
7 Hello World!
8 SUCCESSFULLY EXECUTED System.Void Hello.Hello::SayHello()
9 EXITING System.Void Hello.Hello::SayHello()
10 ENTERING System.Void Hello.Hello::Say(System.String)
11     msg (System.String) = Hey Buffalo how's it going!
12 Hey Buffalo how's it going!
13 SUCCESSFULLY EXECUTED System.Void Hello.Hello::Say(System.String)
14     msg (System.String) = Hey Buffalo how's it going!
15 EXITING System.Void Hello.Hello::Say(System.String)
16     msg (System.String) = Hey Buffalo how's it going!
```

Listing A.7: TraceAspect Output

Line 7 and 12 are the original method output; the rest are the output of the various interception points. Note that line 2, 11, 14 and 16 show the parameter value passed into each method.

A.3 Transaction Database

It is common to work with code that saves information to a database, and it is typical that the data might have to go into different tables. For example, let us look at an online ordering system. When a customer clicks the buy button during checkout, an order has to be created, and each order item has to be created as well. Order and order items usually go into different tables since they have the logical parent-child relationship. Given this, we might have the following code to save the data:

```
1 public void SaveOrder(string customerName, string item1, string item2)
```

```

2  {
3      //create the linq-to-entity context
4      ModelContainer db = new ModelContainer();
5      //create an order obj
6      Order o = new Order();
7      o.CustomerName = customerName;
8      db.Orders.Add(o);
9      db.SaveChanges(); // <--save the order!
10
11     List<string> items = new List<string>();
12     if (!string.IsNullOrEmpty(item1)) items.Add(item1);
13     if (!string.IsNullOrEmpty(item2)) items.Add(item2);
14     items.ForEach(x =>
15     {
16         //create the order item
17         OrderItem item = new OrderItem();
18         item.ItemName = x;
19         item.OrderId = -1; //<-- intentionally failing the save since there is no such order id
20         db.OrderItems.Add(item);
21     });
22
23     db.SaveChanges(); //<-- saving the order items
24 }

```

Listing A.8: SaveOrder Example

This code is for illustration purposes only. It first saves an order and then saves each individual order items. The data goes into different tables. Note that we are using the Linq to Entity ORM here. Technically it handles the transaction if `db.SaveChanges()` at line 9 is removed and is called at the end with one call.

Also note that at line 19, we are intentionally failing the `OrderItem` save. At that point, the `Order` itself has been saved. When `OrderItem` fails, we ended up with incomplete data in the database. In a case like this, it is important that the atomicity rule is enforced. The save operation should either be completely saved or nothing should be saved at all.

Atomicity rule is usually enforced by using transaction. So we can modify the above code like this:

```

1  public void SaveOrderManualTransaction(string customerName, string item1, string item2)
2  {
3      using (TransactionScope scope = new TransactionScope())
4      {
5          //same body of code to save..
6      }
7  }

```

Listing A.9: SaveOrderManualTransaction Example

The save operations are now wrapped in a `TransactionScope` object. If any unforeseeable error happens during any of the `db.SaveChanges()` calls, the whole operation is aborted, so we will not end up with incomplete data in the database.

A scenario like this is a suitable candidate for using aspect. If we need to enforce atomicity elsewhere, we would have to do something similar with `TransactionScope`. We can refactor and extract out the aspect using Buffalo:

```
1 public class RunInTransactionAspect : MethodBoundaryAspect
2 {
3     private TransactionScope scope;
4
5     public override void OnBefore(MethodArgs args)
6     {
7         this.scope = new TransactionScope(TransactionScopeOption.RequiresNew);
8     }
9
10    public override void OnSuccess(MethodArgs args)
11    {
12        this.scope.Complete();
13    }
14
15    public override void OnException(MethodArgs args)
16    {
17        Transaction.Current.Rollback();
18    }
19
20    public override void OnAfter(MethodArgs args)
21    {
22        this.scope.Dispose();
23    }
24 }
```

Listing A.10: RunInTransactionAspect

Then apply this aspect to the original code:

```
1 [RunInTransactionAspect]
2 public void SaveOrder(string customerName, string item1, string item2)
3 {
4     //same body of code to save..
5 }
```

Listing A.11: SaveOrder With Aspect

This will ensure that the method will operate in the transaction in a safe manner. Now we can reuse this aspect by applying it to other operations that need to be transaction scoped.

A.4 Update UI Safely with Thread




Figure A.1: UI Thread Example

This is a common problem faced by Windows developers. Suppose that when the Save button is clicked, a long-running operation is triggered. Here, we stimulate the long running operation by putting the thread to sleep. The label that says Ready is updated to OK when the operation finishes:

```
1 public void SaveData ()
2 {
3     //sleep to simulate long running process
4     Thread.Sleep(10000);
5     lblStatus.Text = "OK";
6 }
```

Listing A.12: Long-running Operation

The window application is single thread application. While the long operation is in progress, the whole application will be unresponsive until the operation finishes. This makes for a very bad user experience.

One remedy is to put the long-running operation in a separate worker thread and update the UI label when the worker thread finishes.

```
1 public void SaveData ()
2 {
3     new Thread(new ThreadStart(() =>
4     {
5         Thread.Sleep(5000);
6         lblStatus.Text = "OK";
7     })).Start();
8 }
```

Listing A.13: Long-running Operation With Thread

However, this would cause an exception to be thrown on line 6. UI elements can only be updated from the thread that created them.

```
1 public void SaveData()
2 {
3     Thread.Sleep(5000);
4     Action action = () => lblStatus.Text = "OK";
5     this.BeginInvoke(action);
6 }
```

Listing A.14: Update UI From Different Thread

Listing A.14 will marshal the request to update the UI back to the UI thread. This ensures that the label will be updated without any threading problem. Since this is a common threading problem when working with UIs, Buffalo can be used to extract the solution out as an aspect:

```
1 public class ThreadSafeAspect : MethodAroundAspect
2 {
3     public override object Invoke(Buffalo.MethodArgs args)
4     {
5         Dispatcher.CurrentDispatcher.Invoke(() => args.Proceed());
6         return null;
7     }
8 }
```

Listing A.15: ThreadSafeAspect

Here, we create a new aspect inheriting from MethodAroundAspect and overriding Invoke. This wraps around the invocation of the method call. Args.Proceed() is Buffalos way of calling back to the original UpdateStatus method:

```
1 public void SaveData()
2 {
3     new Thread(new ThreadStart(() =>
4     {
5         Thread.Sleep(5000);
6         this.UpdateStatus();
7     })).Start();
8 }
9 [ThreadSafeAspect]
10 public void UpdateStatus() { lblStatus.Text = "OK"; }
```

Listing A.16: Apply ThreadSafeAspect

Now when the Save button is clicked, the long-running operation is triggered, and control is immediately returned. When the operation finishes, the label is updated without any threading issue.

Note that, currently, `MethodAroundAspect` supports only method-level attribution. This limits its reusability. It will be much more useful when it supports three levels as `MethodBoundaryAspect` does.

A.5 Integrate With MS-Build System

Buffalo can be integrated with MS-Build, so weaving can be invoked automatically when a project is compiled from the Visual Studio IDE. Note that the following instructions are just the bare minimum to get this working; many of the bells and whistles are omitted.

MS-Build is integrated with Visual Studio IDE via a configuration file. For example, a C# project has the associated `.csproj`; if the project file is opened in a text editor you will see a line that references a different configuration file: `Microsoft.CSharp.targets`. This file, in turn, references `Microsoft.Common.targets`.

Each .NET version has its own `Microsoft.Common.targets` file. Depending on the version you are using, open up this file in a text editor. For example, for .NET 4.0, this file is located in `C:\Windows\Microsoft.NET\Framework\v4.0.30319\Microsoft.Common.targets`.

Under the Compile section, around line #2013, add the line to import the `Buffalo.targets` file as shown in Figure A.2



Figure A.2: Adding Buffalo.targets

If you open Buffalo.targets in a text editor, it contains the following context:

```

1 <Project xmlns="http://schemas.microsoft.com/developer/msbuild/2003">
2   <PropertyGroup>
3     <CompileDependsOn>
4       $(CompileDependsOn);
5       Buffalo
6     </CompileDependsOn>
7   </PropertyGroup>
8
9   <Target Name="Buffalo">
10    <Message Text="Hello Buffalo! @(IntermediateAssembly)"/>
11    <Exec Command="&quot;C:\Buffalo\BuffaloAOP.exe&quot; &quot;%(IntermediateAssembly)&quot;"/>
12  </Target>
13 </Project>

```

Listing A.17: Buffalo.targets

This is how Buffalo gets hooked into MS-Build. What this means is that when the user compiles a project, everything defined in the CompileDependsOn property group will be performed first. Then, a new target named "Buffalo" will be called immediately, which will invoke the BuffaloAOP via the Exec Command. Note that for the Exec Command, a complete path to BuffaloAOP.exe must be provided, including the encoded quotation marks as shown.

Make sure to save all the changes.

Now every time a C# project is compiled, Buffalo will be invoked automatically to perform the weaving.